

# homework3\_stub

May 1, 2025

## 0.1 Homework 3: Symbolic Music Generation Using Markov Chains

### Before starting the homework:

Please run `pip install miditok` to install the [MiDiTok](#) package, which simplifies MIDI file processing by making note and beat extraction more straightforward.

You're also welcome to experiment with other MIDI processing libraries such as [mido](#), [pretty\\_midi](#) and [miditoolkit](#). However, with these libraries, you'll need to handle MIDI quantization yourself, for example, converting note-on/note-off events into beat positions and durations.

```
[ ]: # run this command to install MiDiTok
# ! pip install miditok
```

Collecting miditok

Downloading miditok-3.0.5.post1-py3-none-any.whl.metadata (10 kB)

Collecting huggingface-hub>=0.16.4 (from miditok)

Downloading huggingface-hub-0.30.2-py3-none-any.whl.metadata (13 kB)

Requirement already satisfied: numpy>=1.19 in /opt/homebrew/lib/python3.10/site-packages (from miditok) (1.26.3)

Collecting symusic>=0.5.0 (from miditok)

Downloading symusic-0.5.7-cp310-cp310-macosx\_11\_0\_arm64.whl.metadata (8.7 kB)

Collecting tokenizers>=0.13.0 (from miditok)

Downloading tokenizers-0.21.1-cp39-abi3-macosx\_11\_0\_arm64.whl.metadata (6.8 kB)

Requirement already satisfied: tqdm in /opt/homebrew/lib/python3.10/site-packages (from miditok) (4.66.2)

Requirement already satisfied: filelock in /opt/homebrew/lib/python3.10/site-packages (from huggingface-hub>=0.16.4->miditok) (3.13.1)

Requirement already satisfied: fsspec>=2023.5.0 in /opt/homebrew/lib/python3.10/site-packages (from huggingface-hub>=0.16.4->miditok) (2024.2.0)

Requirement already satisfied: packaging>=20.9 in /Users/ryoandrewonozuka/Library/Python/3.10/lib/python/site-packages (from huggingface-hub>=0.16.4->miditok) (23.2)

Requirement already satisfied: pyyaml>=5.1 in /opt/homebrew/lib/python3.10/site-packages (from huggingface-hub>=0.16.4->miditok) (6.0.1)

Requirement already satisfied: requests in /opt/homebrew/lib/python3.10/site-packages (from huggingface-hub>=0.16.4->miditok) (2.31.0)

Requirement already satisfied: typing-extensions>=3.7.4.3 in

```

/opt/homebrew/lib/python3.10/site-packages (from huggingface-
hub>=0.16.4->miditok) (4.10.0)
Collecting pySmartDL (from symusic>=0.5.0->miditok)
  Downloading pySmartDL-1.3.4-py3-none-any.whl.metadata (2.8 kB)
Requirement already satisfied: platformdirs in
/Users/ryoandrewonozuka/Library/Python/3.10/lib/python/site-packages (from
symusic>=0.5.0->miditok) (4.1.0)
Requirement already satisfied: charset-normalizer<4,>=2 in
/opt/homebrew/lib/python3.10/site-packages (from requests->huggingface-
hub>=0.16.4->miditok) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in
/opt/homebrew/lib/python3.10/site-packages (from requests->huggingface-
hub>=0.16.4->miditok) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/opt/homebrew/lib/python3.10/site-packages (from requests->huggingface-
hub>=0.16.4->miditok) (2.2.1)
Requirement already satisfied: certifi>=2017.4.17 in
/opt/homebrew/lib/python3.10/site-packages (from requests->huggingface-
hub>=0.16.4->miditok) (2024.2.2)
Downloading miditok-3.0.5.post1-py3-none-any.whl (158 kB)
Downloading huggingface_hub-0.30.2-py3-none-any.whl (481 kB)
Downloading symusic-0.5.7-cp310-cp310-macosx_11_0_arm64.whl (2.6 MB)
      2.6/2.6 MB
25.4 MB/s eta 0:00:00
Downloading tokenizers-0.21.1-cp39-abi3-macosx_11_0_arm64.whl (2.7 MB)
      2.7/2.7 MB
17.1 MB/s eta 0:00:00
Downloading pySmartDL-1.3.4-py3-none-any.whl (20 kB)
Installing collected packages: pySmartDL, symusic, huggingface-hub, tokenizers,
miditok
Successfully installed huggingface-hub-0.30.2 miditok-3.0.5.post1
pySmartDL-1.3.4 symusic-0.5.7 tokenizers-0.21.1

[notice] A new release of pip is
available: 25.0 -> 25.1
[notice] To update, run:
python3.10 -m pip install --upgrade pip

```

```

[2]: # import required packages
import random
from glob import glob
from collections import defaultdict

import numpy as np
from numpy.random import choice

from symusic import Score

```

```
from miditok import REMI, TokenizerConfig
from midiutil import MIDIFile
```

```
[ ]: # You can change the random seed but try to keep your results deterministic!
# If I need to make changes to the autograder it'll require rerunning your code,
# so it should ideally generate the same results each time.
random.seed(42)
```

### 0.1.1 Load music dataset

We will use a subset of the [PDMX dataset](#).

Please find the link in the homework spec.

All pieces are monophonic music (i.e. one melody line) in 4/4 time signature.

```
[ ]: midi_files = glob('PDMX_subset/*.mid')
len(midi_files)
```

### 0.1.2 Train a tokenizer with the REMI method in MidiTok

```
[ ]: config = TokenizerConfig(num_velocities=1, use_chords=False, use_programs=False)
tokenizer = REMI(config)
tokenizer.train(vocab_size=1000, files_paths=midi_files)
```

### 0.1.3 Use the trained tokenizer to get tokens for each midi file

In REMI representation, each note will be represented with four tokens: Position, Pitch, Velocity, Duration, e.g. ('Position\_28', 'Pitch\_74', 'Velocity\_127', 'Duration\_0.4.8'); a Bar\_None token indicates the beginning of a new bar.

```
[ ]: # e.g.:
midi = Score(midi_files[0])
tokens = tokenizer(midi)[0].tokens
tokens[:10]
```

1. Write a function to extract note pitch events from a midi file; and another extract all note pitch events from the dataset and output a dictionary that maps note pitch events to the number of times they occur in the files. (e.g. {60: 120, 61: 58, ...}).

`note_extraction()` - **Input:** a midi file

- **Output:** a list of note pitch events (e.g. [60, 62, 61, ...])

`note_frequency()` - **Input:** all midi files `midi_files`

- **Output:** a dictionary that maps note pitch events to the number of times they occur, e.g. {60: 120, 61: 58, ...}

```
[ ]: def note_extraction(midi_file):
    # Q1a: Your code goes here
    pass
```

```
[ ]: def note_frequency(midi_files):
    # Q1b: Your code goes here
    pass
```

2. Write a function to normalize the above dictionary to produce probability scores (e.g. {60: 0.13, 61: 0.065, ...})

`note_unigram_probability()` - **Input:** all midi files `midi_files`

- **Output:** a dictionary that maps note pitch events to probabilities, e.g. {60: 0.13, 61: 0.06, ...}

```
[ ]: def note_unigram_probability(midi_files):
    note_counts = note_frequency(midi_files)
    unigramProbabilities = {}

    # Q2: Your code goes here
    # ...

    return unigramProbabilities
```

3. Generate a table of pairwise probabilities containing  $p(\text{next\_note} \mid \text{previous\_note})$  values for the dataset; write a function that randomly generates the next note based on the previous note based on this distribution.

`note_bigram_probability()` - **Input:** all midi files `midi_files`

- **Output:** two dictionaries:
  - `bigramTransitions`: key: `previous_note`, value: a list of `next_note`, e.g. {60:[62, 64, ..], 62:[60, 64, ..], ...} (i.e., this is a list of every other note that occurred after note 60, every note that occurred after note 62, etc.)
  - `bigramTransitionProbabilities`: key:`previous_note`, value: a list of probabilities for `next_note` in the same order of `bigramTransitions`, e.g. {60:[0.3, 0.4, ..], 62:[0.2, 0.1, ..], ...} (i.e., you are converting the values above to probabilities)

`sample_next_note()` - **Input:** a note

- **Output:** next note sampled from pairwise probabilities

```
[ ]: def note_bigram_probability(midi_files):
    bigramTransitions = defaultdict(list)
    bigramTransitionProbabilities = defaultdict(list)

    # Q3a: Your code goes here
    # ...

    return bigramTransitions, bigramTransitionProbabilities
```

```
[ ]: def sample_next_note(note):
    # Q3b: Your code goes here
```

```
pass
```

4. Write a function to calculate the perplexity of your model on a midi file.

The perplexity of a model is defined as

$$\exp(-\frac{1}{N} \sum_{i=1}^N \log(p(w_i|w_{i-1})))$$

where  $p(w_1|w_0) = p(w_1)$ ,  $p(w_i|w_{i-1}) (i > 1)$  refers to the pairwise probability  $p(\text{next\_note} | \text{previous\_note})$ .

`note_bigram_perplexity()` - **Input:** a midi file

- **Output:** perplexity value

```
[ ]: def note_bigram_perplexity(midi_file):
    unigramProbabilities = note_unigram_probability(midi_files)
    bigramTransitions, bigramTransitionProbabilities = \
    ↪note_bigram_probability(midi_files)

    # Q4: Your code goes here
    # Can use regular numpy.log (i.e., natural logarithm)
```

5. Implement a second-order Markov chain, i.e., one which estimates  $p(\text{next\_note} | \text{next\_previous\_note}, \text{previous\_note})$ ; write a function to compute the perplexity of this new model on a midi file.

The perplexity of this model is defined as

$$\exp(-\frac{1}{N} \sum_{i=1}^N \log(p(w_i|w_{i-2}, w_{i-1})))$$

where  $p(w_1|w_{-1}, w_0) = p(w_1)$ ,  $p(w_2|w_0, w_1) = p(w_2|w_1)$ ,  $p(w_i|w_{i-2}, w_{i-1}) (i > 2)$  refers to the probability  $p(\text{next\_note} | \text{next\_previous\_note}, \text{previous\_note})$ .

`note_trigram_probability()` - **Input:** all midi files `midi_files`

- **Output:** two dictionaries:

- `trigramTransitions`: key -  $(\text{next\_previous\_note}, \text{previous\_note})$ , value - a list of `next\_note`, e.g.  $\{(60, 62):[64, 66, \dots], (60, 64):[60, 64, \dots], \dots\}$
- `trigramTransitionProbabilities`: key:  $(\text{next\_previous\_note}, \text{previous\_note})$ , value: a list of probabilities for `next\_note` in the same order of `trigramTransitions`, e.g.  $\{(60, 62):[0.2, 0.2, \dots], (60, 64):[0.4, 0.1, \dots], \dots\}$

`note_trigram_perplexity()` - **Input:** a midi file

- **Output:** perplexity value

```
[ ]: def note_trigram_probability(midi_files):
    trigramTransitions = defaultdict(list)
    trigramTransitionProbabilities = defaultdict(list)

    # Q5a: Your code goes here
    # ...
```

```
return trigramTransitions, trigramTransitionProbabilities
```

```
[ ]: def note_trigram_perplexity(midi_file):
    unigramProbabilities = note_unigram_probability(midi_files)
    bigramTransitions, bigramTransitionProbabilities = \
    ↪note_bigram_probability(midi_files)
    trigramTransitions, trigramTransitionProbabilities = \
    ↪note_trigram_probability(midi_files)

    # Q5b: Your code goes here
```

6. Our model currently doesn't have any knowledge of beats. Write a function that extracts beat lengths and outputs a list of [(beat position; beat length)] values.

Recall that each note will be encoded as **Position**, **Pitch**, **Velocity**, **Duration** using REMI. Please keep the **Position** value for beat position, and convert **Duration** to beat length using provided lookup table `duration2length` (see below).

For example, for a note represented by four tokens ('Position\_24', 'Pitch\_72', 'Velocity\_127', 'Duration\_0.4.8'), the extracted (beat position; beat length) value is (24, 4).

As a result, we will obtain a list like [(0,8),(8,16),(24,4),(28,4),(0,4)...], where the next beat position is the previous beat position + the beat length. As we divide each bar into 32 positions by default, when reaching the end of a bar (i.e.  $28 + 4 = 32$  in the case of (28, 4)), the beat position reset to 0.

```
[ ]: duration2length = {
    '0.2.8': 2, # sixteenth note, 0.25 beat in 4/4 time signature
    '0.4.8': 4, # eighth note, 0.5 beat in 4/4 time signature
    '1.0.8': 8, # quarter note, 1 beat in 4/4 time signature
    '2.0.8': 16, # half note, 2 beats in 4/4 time signature
    '4.0.4': 32, # whole note, 4 beats in 4/4 time signature
}
```

`beat_extraction()` - **Input:** a midi file

- **Output:** a list of (beat position; beat length) values

```
[ ]: def beat_extraction(midi_file):
    # Q6: Your code goes here
    pass
```

7. Implement a Markov chain that computes  $p(\text{beat\_length} \mid \text{previous\_beat\_length})$  based on the above function.

`beat_bigram_probability()` - **Input:** all midi files `midi_files`

- **Output:** two dictionaries:

- `bigramBeatTransitions`: key: `previous_beat_length`, value: a list of `beat_length`, e.g. `{4:[8, 2, ..], 8:[8, 4, ..], ...}`
- `bigramBeatTransitionProbabilities`: key - `previous_beat_length`, value - a list of probabilities for `beat_length` in the same order of `bigramBeatTransitions`, e.g. `{4:[0.3, 0.2, ..], 8:[0.4, 0.4, ..], ...}`

```
[ ]: def beat_bigram_probability(midi_files):
    bigramBeatTransitions = defaultdict(list)
    bigramBeatTransitionProbabilities = defaultdict(list)

    # Q7: Your code goes here
    # ...

    return bigramBeatTransitions, bigramBeatTransitionProbabilities
```

8. Implement a function to compute  $p(\text{beat length} \mid \text{beat position})$ , and compute the perplexity of your models from Q7 and Q8. For both models, we only consider the probabilities of predicting the sequence of **beat lengths**.

`beat_pos_bigram_probability()` - **Input**: all midi files `midi_files`

- **Output**: two dictionaries:
  - `bigramBeatPosTransitions`: key - `beat_position`, value - a list of `beat_length`
  - `bigramBeatPosTransitionProbabilities`: key - `beat_position`, value - a list of probabilities for `beat_length` in the same order of `bigramBeatPosTransitions`

`beat_bigram_perplexity()` - **Input**: a midi file

- **Output**: two perplexity values correspond to the models in Q7 and Q8, respectively

```
[ ]: def beat_pos_bigram_probability(midi_files):
    bigramBeatPosTransitions = defaultdict(list)
    bigramBeatPosTransitionProbabilities = defaultdict(list)

    # Q8a: Your code goes here
    # ...

    return bigramBeatPosTransitions, bigramBeatPosTransitionProbabilities
```

```
[ ]: def beat_bigram_perplexity(midi_file):
    bigramBeatTransitions, bigramBeatTransitionProbabilities =
    ↪ beat_bigram_probability(midi_files)
    bigramBeatPosTransitions, bigramBeatPosTransitionProbabilities =
    ↪ beat_pos_bigram_probability(midi_files)
    # Q8b: Your code goes here
    # Hint: one more probability function needs to be computed

    # perplexity for Q7
```

```

perplexity_Q7 = None

# perplexity for Q8
perplexity_Q8 = None

return perplexity_Q7, perplexity_Q8

```

- Implement a Markov chain that computes  $p(\text{beat\_length} \mid \text{previous\_beat\_length}, \text{beat\_position})$ , and report its perplexity.

`beat_trigram_probability()` - **Input:** all midi files `midi_files`

- Output:** two dictionaries:
  - `trigramBeatTransitions`: key: `(previous_beat_length, beat_position)`, value: a list of `beat_length`
  - `trigramBeatTransitionProbabilities`: key: `(previous_beat_length, beat_position)`, value: a list of probabilities for `beat_length` in the same order of `trigramBeatTransitions`

`beat_trigram_perplexity()` - **Input:** a midi file

- Output:** perplexity value

```

[ ]: def beat_trigram_probability(midi_files):
    trigramBeatTransitions = defaultdict(list)
    trigramBeatTransitionProbabilities = defaultdict(list)

    # Q9a: Your code goes here
    # ...

    return trigramBeatTransitions, trigramBeatTransitionProbabilities

```

```

[ ]: def beat_trigram_perplexity(midi_file):
    bigramBeatPosTransitions, bigramBeatPosTransitionProbabilities = \
    ↪ beat_pos_bigram_probability(midi_files)
    trigramBeatTransitions, trigramBeatTransitionProbabilities = \
    ↪ beat_trigram_probability(midi_files)
    # Q9b: Your code goes here

```

- Use the model from Q5 to generate  $N$  notes, and the model from Q8 to generate beat lengths for each note. Save the generated music as a midi file (see code from workbook1) as `q10.mid`. Remember to reset the beat position to 0 when reaching the end of a bar.

`music_generate` - **Input:** target length, e.g. 500

- Output:** a midi file `q10.mid`

Note: the duration of one beat in `MIDIUtil` is 1, while in `MidiTok` is 8. Divide beat length by 8 if you use methods in `MIDIUtil` to save midi files.



```
[ ]: def music_generate(length):
    # sample notes
    unigramProbabilities = note_unigram_probability(midi_files)
    bigramTransitions, bigramTransitionProbabilities = \
    ↪note_bigram_probability(midi_files)
    trigramTransitions, trigramTransitionProbabilities = \
    ↪note_trigram_probability(midi_files)

    # Q10: Your code goes here ...
    sampled_notes = []

    # sample beats
    sampled_beats = []

    # save the generated music as a midi file
```