



CSE 167 (WI25)

Computer Graphics: Graphics Pipeline

Albert Chern

Overview

- An **algorithm** for drawing picture
- What is **rasterization**? What is **ray tracing**?
- What is **fragment**? What is **buffer**? What is **shader**?
- How do you **describe a shape**?
- What is **GPU**?

The Main Algorithm

- Main algorithm
- Rasterization pipeline
- GPU
- Summary

Big question

How to draw pictures algorithmically?

Big question

How to draw pictures algorithmically?

- Prescribe/input:
- Output:

Big question

How to draw pictures algorithmically?

- Prescribe/input:
 - ▶ Geometries (triangle mesh) in the scene
 - ▶ Color, material,... of each geometry
 - ▶ Light
 - ▶ Camera
- Output:
 - ▶ Color per pixel in the screen



Forza Motorsport 7
(Turn 10 Studios, Microsoft.)

Main algorithm

- Prescribe/input:
 - ▶ Geometries (triangle mesh) in the scene
 - ▶ Color, material,... of each geometry
 - ▶ Light
 - ▶ Camera

Step 1

Determine which triangle corresponds to which pixel


Step 2

Color each pixel according to lighting/material/orientation etc.

- Output:
 - ▶ Color per pixel in the screen

From 3D to 2D

3D scene





center of projection

screen

these are
“camera”


From 3D to 2D



- Working out which geometries correspond to which pixels can be tedious

From 3D to 2D

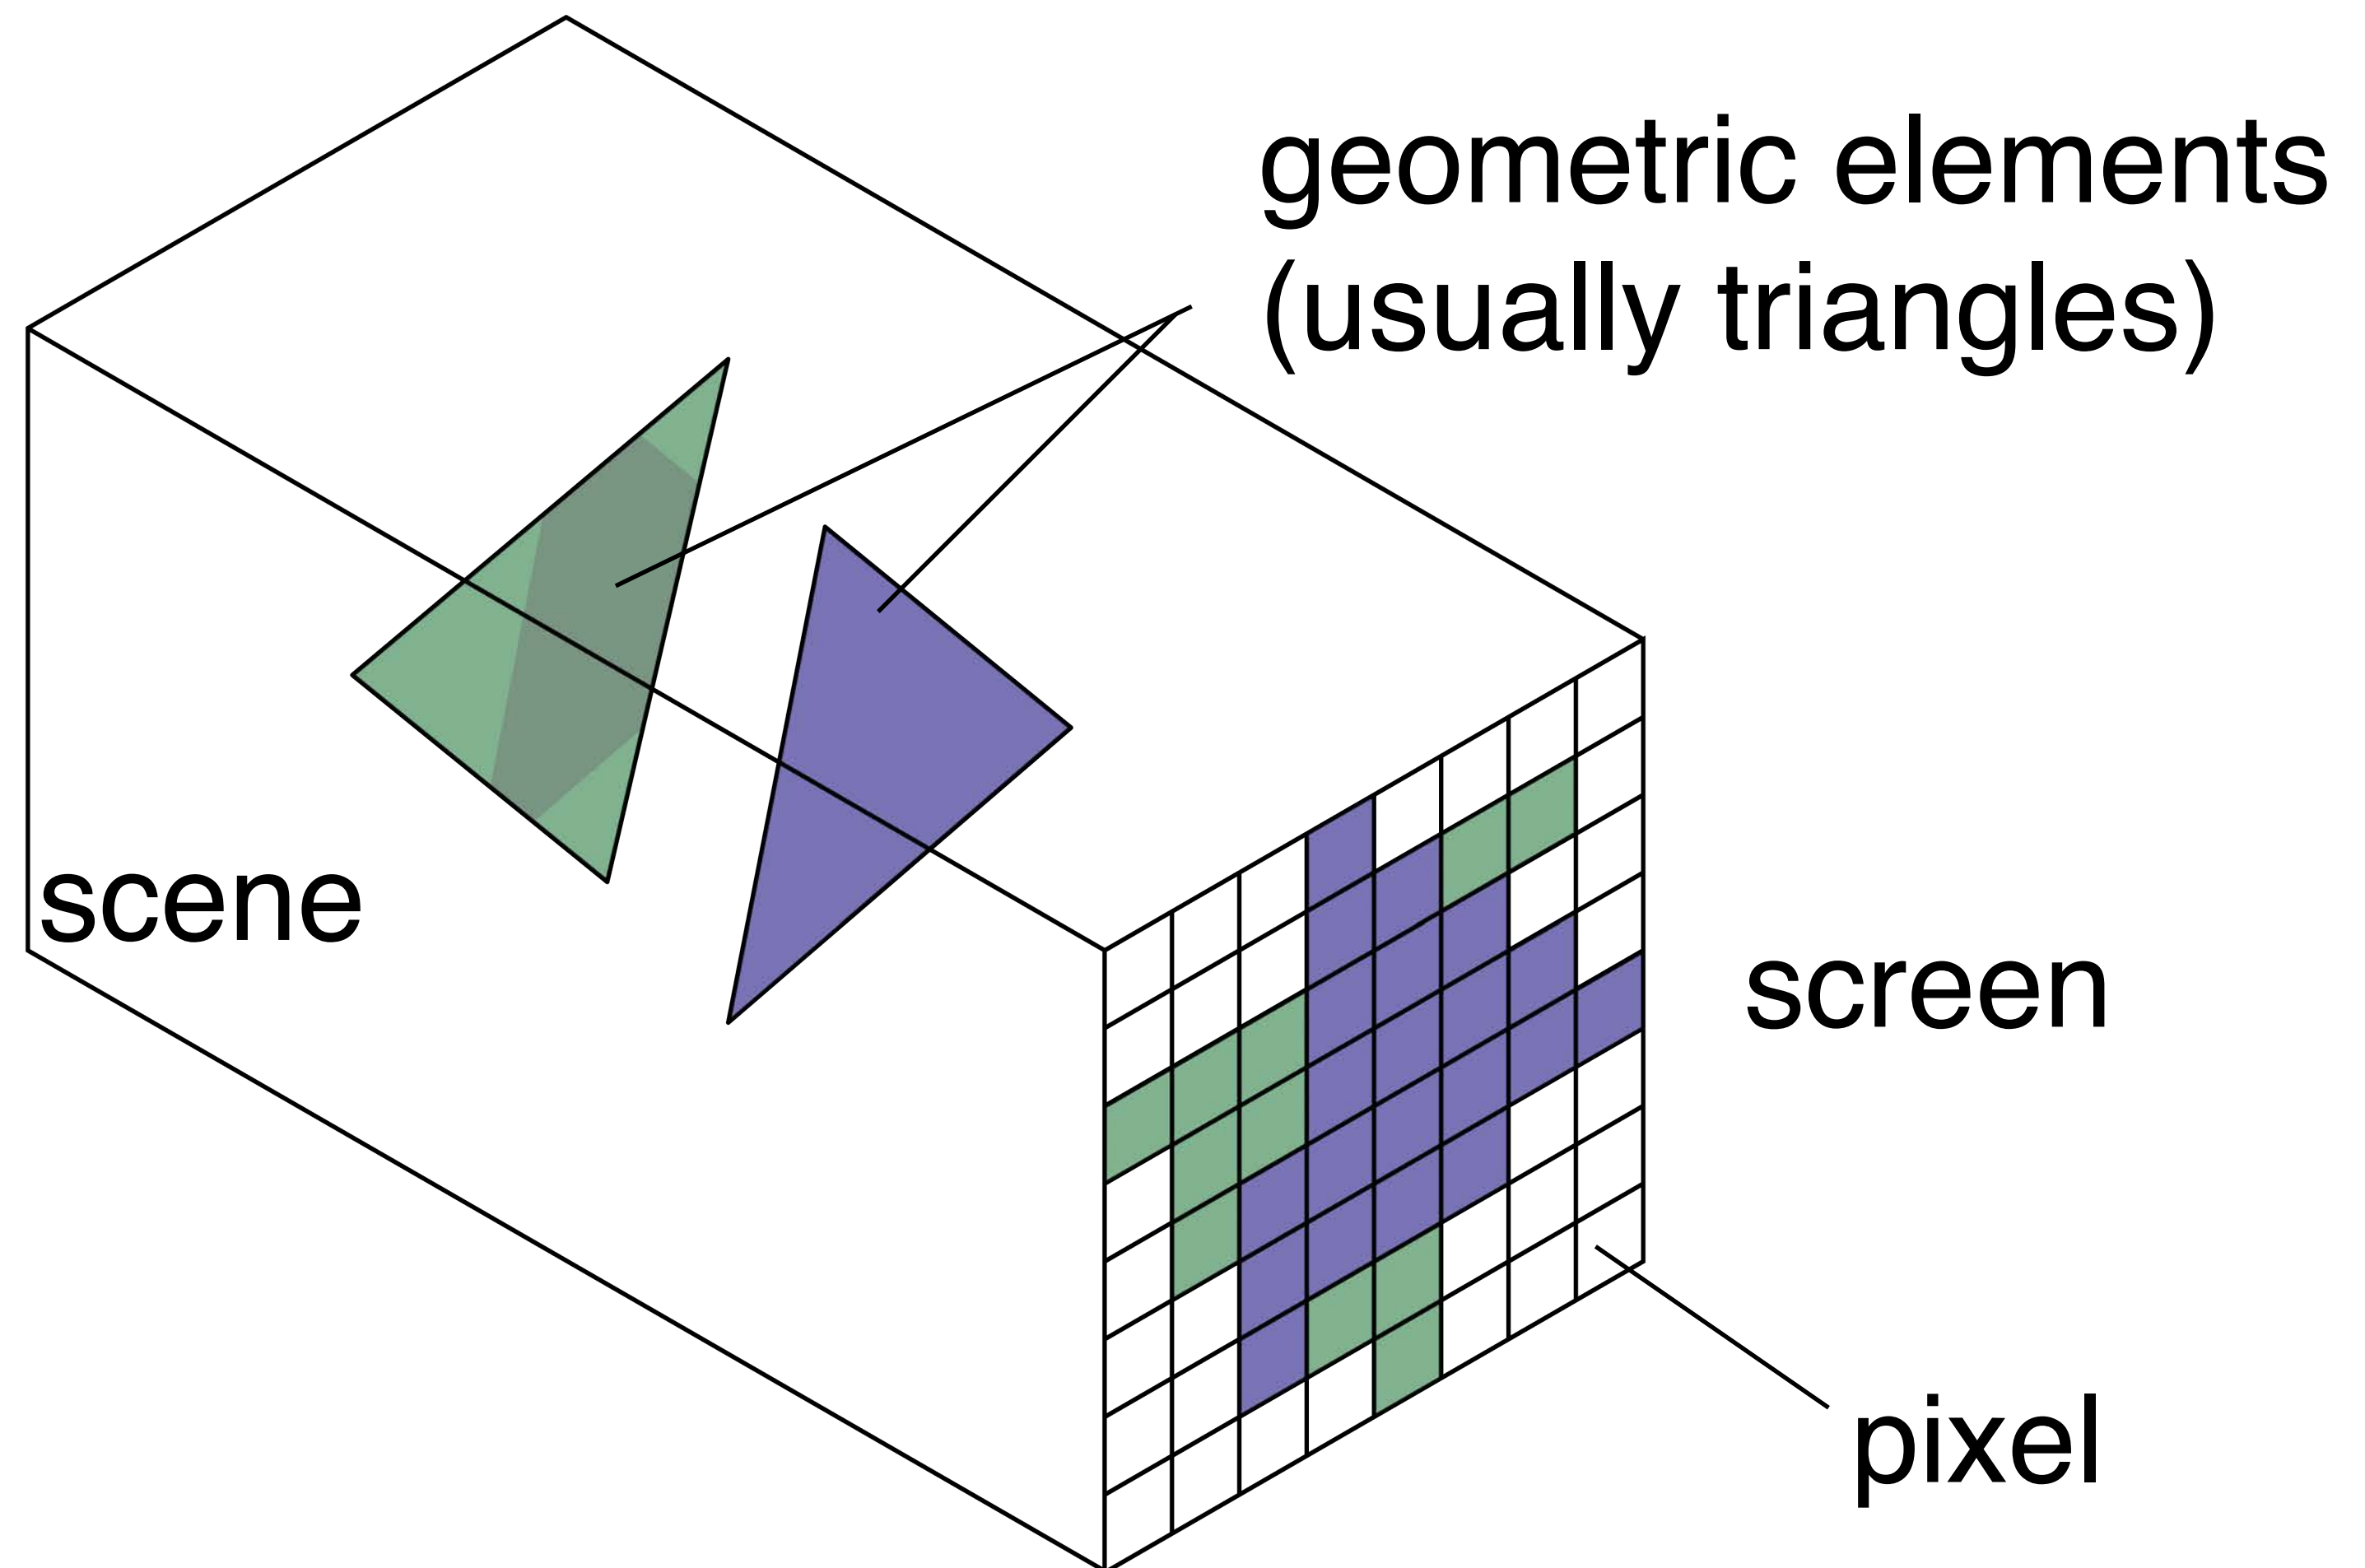

- In week 4, we will learn that there is always a transformation that normalizes the frustum into a standard box



- The question remains. How do we turn triangles to pixels?


Scene-Screen Incidence Relation

- Determine which triangles of the scene is incident to which pixels of the screen



Scene-Screen Incidence Relation

- Determine which triangles of the scene is incident to which pixels of the screen



Definition


A *fragment* is a triangle-pixel pair so that the pixel is incident to (covered by) the triangle.

- Each fragment knows which triangle it comes from
- Each fragment knows which pixel to land on

Scene-Screen Incidence Relation

We are constructing a mapping between scene and screen.

There are 2 main ways for traversing and listing out all the fragments




Scene-Screen Incidence Relation

Method 1

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

Method 2

```
for each pixel in screen
  for each triangle in scene
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```



Scene-Screen Incidence Relation

Method 1


```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

Method 2

```
for each pixel in screen
  for each triangle in scene
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

We will learn this math problem in Week5:

What is the criterion that the pixel center is in the triangle?




Scene-Screen Incidence Relation

Method 1

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

Method 2

```
for each pixel in screen
  for each triangle in scene
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```




Scene-Screen Incidence Relation

Rasterization

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

Ray casting/tracing

```
for each pixel in screen
  for each triangle in scene
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```




Rasterization v.s. Ray tracing

Rasterization

- Easy to speed up
- Hardcoded in hardware
- Real-time rendering
- No photorealism

Ray casting/tracing

- Hard to speed up
(need nontrivial data structure)
- Hardware available recently
- Recursive ray tracing => photorealism



Rasterization v.s. Ray tracing

Rasterization

- Easy to speed up
- Hardcoded in hardware
- Real-time rendering
- No photorealism



Ray casting/tracing

- Hard to speed up
(need nontrivial data structure)
- Hardware available recently
- Recursive ray tracing => photorealism



Rasterization v.s. Ray tracing

Rasterization

- Easy to speed up
- Hardcoded in hardware
- Real-time rendering
- No photorealism

First few weeks

- ▶ Runs on GPU
- ▶ OpenGL API



Ray casting/tracing

- Hard to speed up
(need nontrivial data structure)
- Hardware available recently
- Recursive ray tracing => photorealism



Rasterization-based Graphics Pipeline

- Main algorithm
- Rasterization pipeline
- GPU
- Summary

Rasterization

Back to the main algorithm

Raster graphics pipeline

rasterizer


```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```


Raster graphics pipeline

Pseudocode


rasterizer

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```




Data in
allocated
memory
a.k.a. **buffer**

list of
triangles




list of
fragments



Raster graphics pipeline

rasterizer


```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```



Raster graphics pipeline

rasterizer

```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```




Raster graphics pipeline


rasterizer

```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```

list of
triangles



list of
fragments



Raster graphics pipeline


rasterizer

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```


fragment shader

```
for each fragment
  ...
  compute color
  ...
  assign frag_color
end for
```


list of
triangles



list of
fragments



list of
fragments



Raster graphics pipeline

rasterizer

```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```


fragment shader

```
for each fragment
  ...
  compute color
  ...
  assign frag_color
end for
```


resolve depth

```
for each pixel
  keep only the
  top fragment
end for
```


list of
fragments



list of
fragments



color per
pixel



Raster graphics pipeline

izer

... in scene
... in screen
... pair **if** the
... occupies that

fragment shader


```
for each fragment  
...  
compute color  
...  
assign frag_color  
end for
```

resolve depth


```
for each pixel  
keep only the  
top fragment  
end for
```

Show on screen
or
Save as image


list of
fragments



list of
fragments



color per
pixel



Frame buffer

Raster graphics pipeline

rasterizer


```
for each triangle in scene
  for each pixel in screen
    output the pair if the
      triangle occupies that
      pixel.
  end for
end for
```

fragmenter


```
for each fragment
  ...
  compute
  ...
  assign
end for
```

How do we organize
the data that
describes a
triangle mesh?

list of
triangles




list of
fragments




Description of triangle mesh

- Every shape (e.g. square) is partitioned into triangles mesh
- Triangle mesh consists of
 - ▶ Vertices
 - ▶ Connectivity



Description of triangle mesh

- Every shape (e.g. square) is partitioned into triangles mesh
- Triangle mesh consists of
 - ▶ Vertices
 - Position
 - Other attributes like color/material
 - ▶ Connectivity
 - Pointers to 3 vertices per triangle



Description of triangle mesh

- Every shape (e.g. square) is partitioned into triangles mesh


- Triangle mesh consists of
 - ▶ Vertices we call the location storing these data the **vertex buffer**

- Position
 - Other attributes like color/material


- ▶ Connectivity

- Pointers to 3 vertices per triangle

we call the location storing these data the **index buffer**



Description of triangle mesh



Description of triangle mesh


- Give an arbitrary order for vertices.
- Give an arbitrary order for triangles.

Vertex buffer

0	-0.5,	-0.5,
1	0.5,	-0.5,
2	0.5,	0.5,
3	-0.5,	0.5

Index buffer

0	0,	1,	3,
1	2,	3,	1



Description of triangle mesh

Every geometry is completely described in a *geometry spreadsheet* (and an instruction how to parse the vertex buffer)

Geometry spreadsheet (a.k.a. vertex array object (VAO))

Vertex buffer(s)

attribute

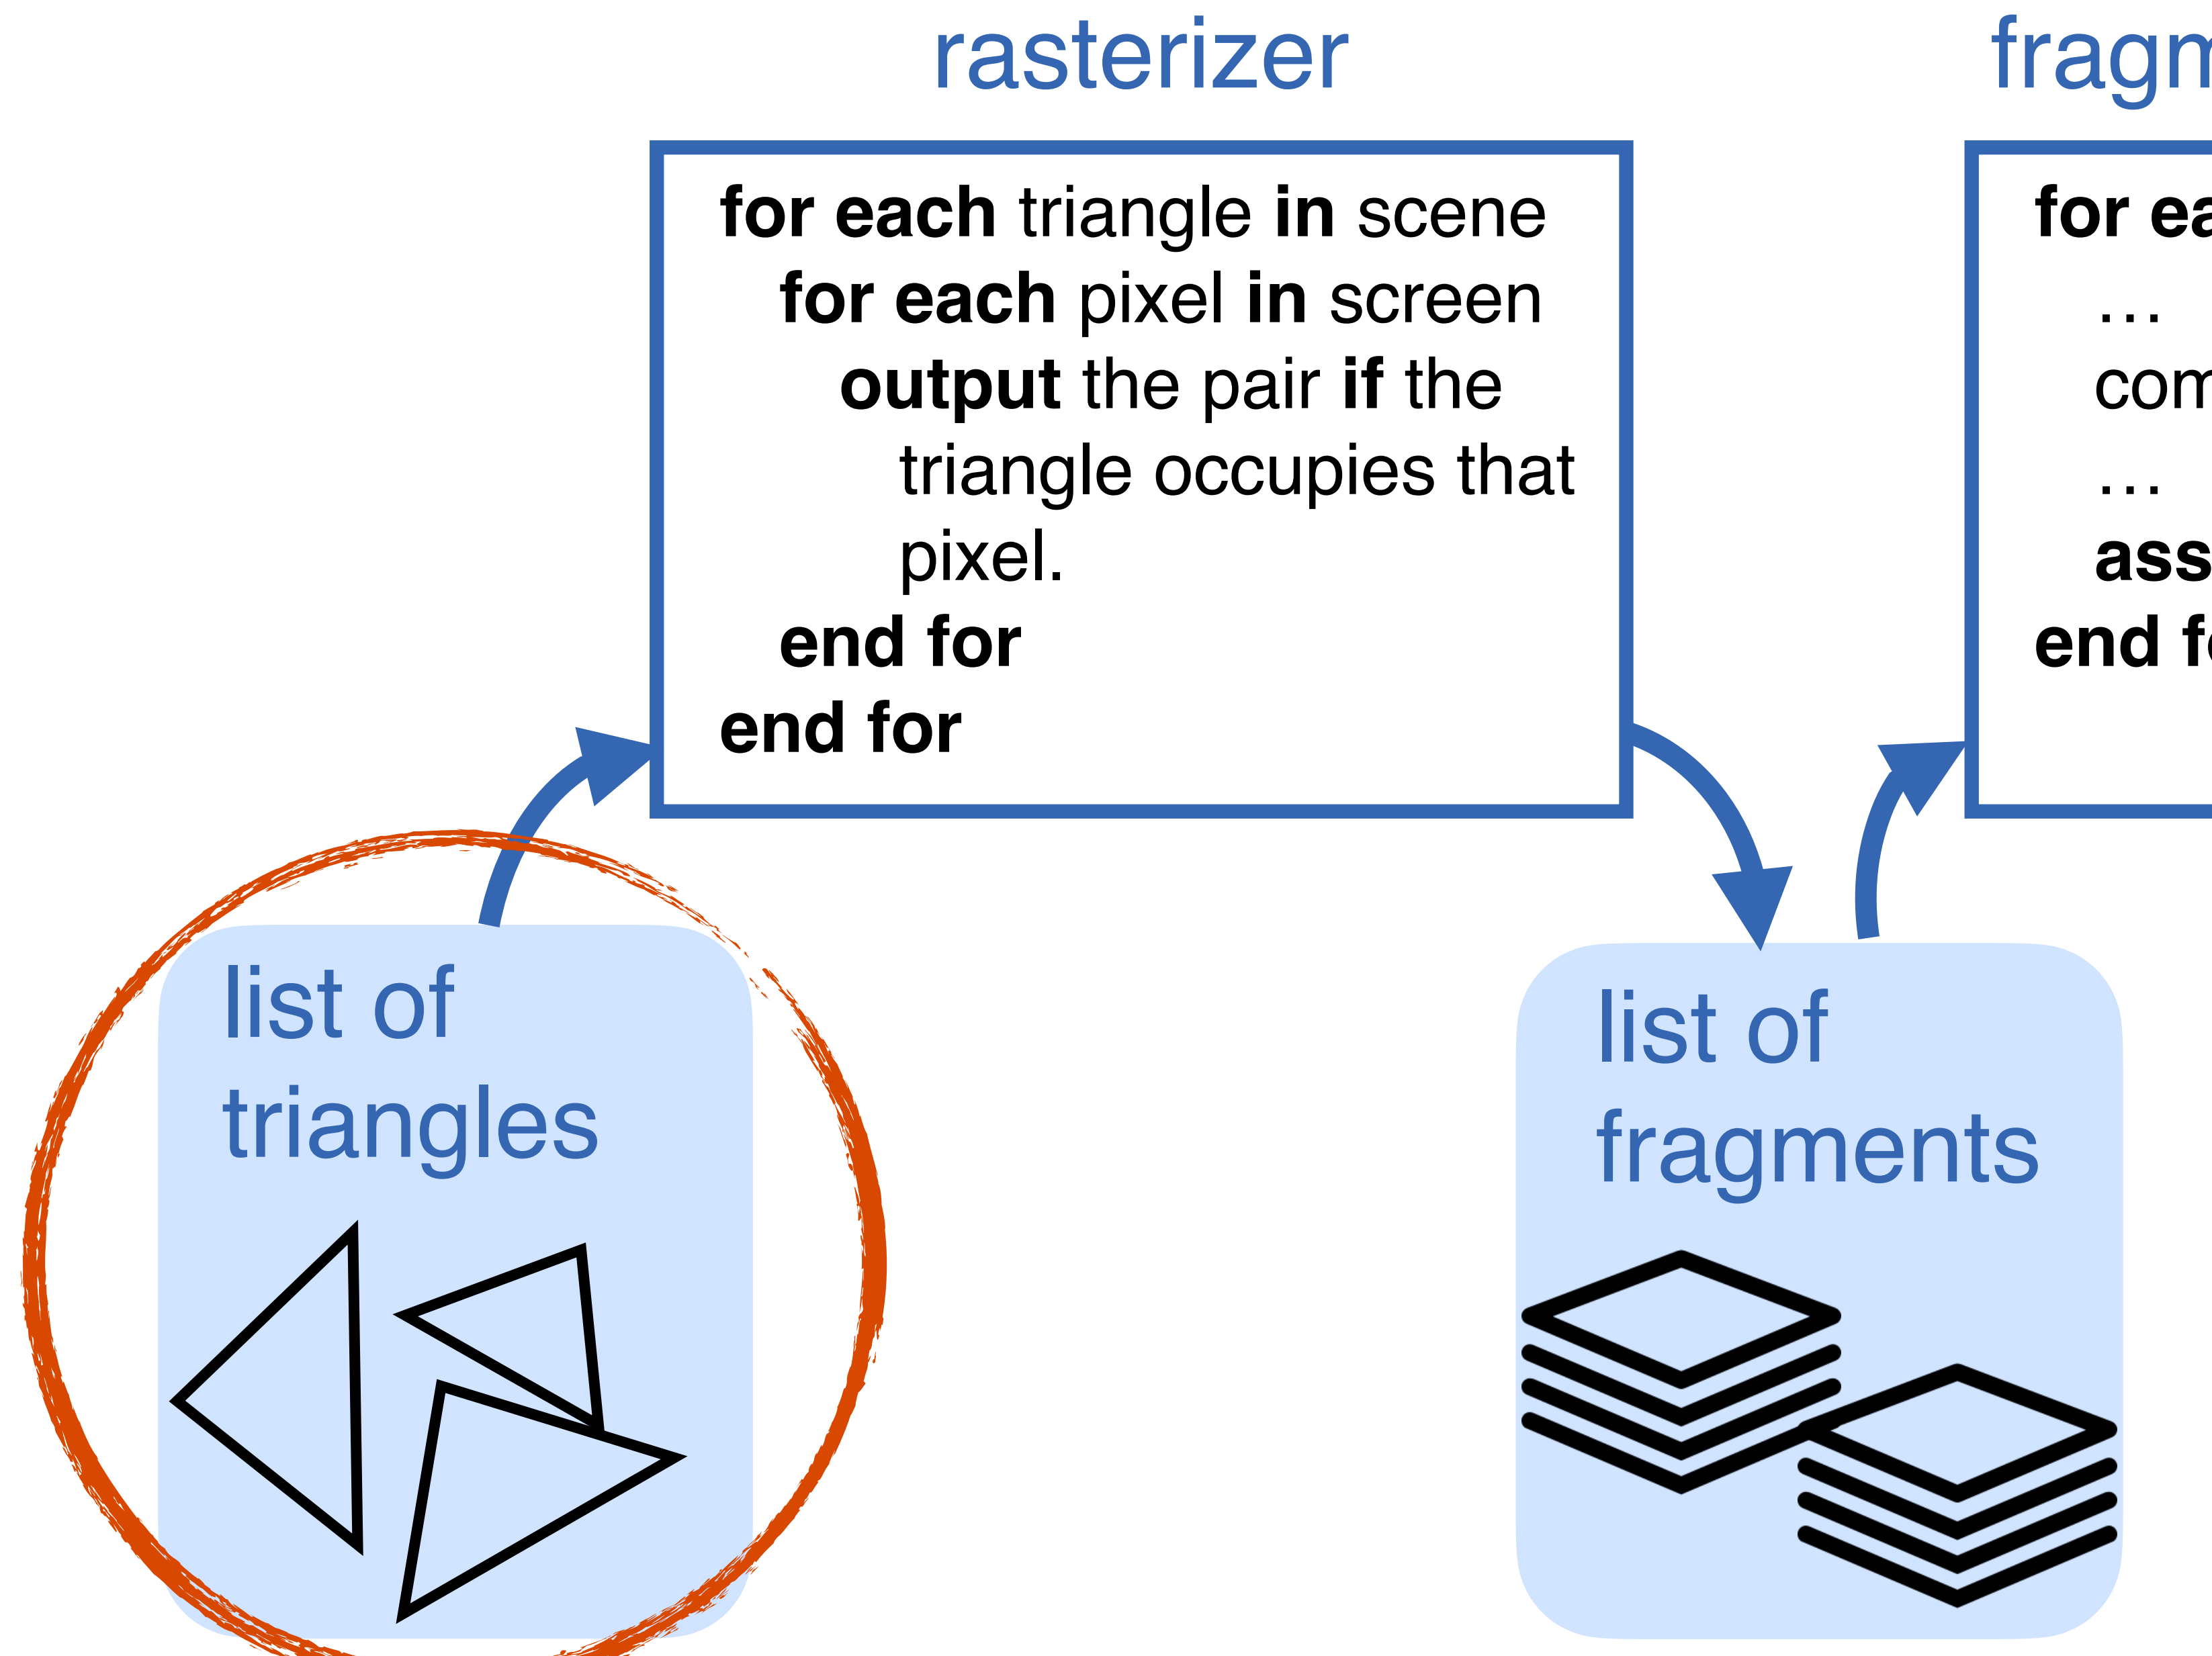
vertex	-0.5	-0.5	0.0	1.0	2.0	0.0
	0.5	-0.5	0.0	1.0	1.0	1.0
	0.5	0.5	0.0	1.0	1.5	-1.5
	-0.5	0.5	0.0	1.0	1.0	1.0
	0.0	0.0	1.0	1.0	0.5	0.5
	0.0	0.0	-1.0	1.0	0.5	0.5
	-0.5	-0.5	0.0	1.0	2.0	0.0
	0.5	-0.5	0.0	1.0	1.0	1.0
	0.5	0.5	0.0	1.0	1.5	-1.5
	-0.5	0.5	0.0	1.0	1.0	1.0
	0.0	0.0	1.0	1.0	0.5	0.5
	0.0	0.0	-1.0	1.0	0.5	0.5

Index buffer

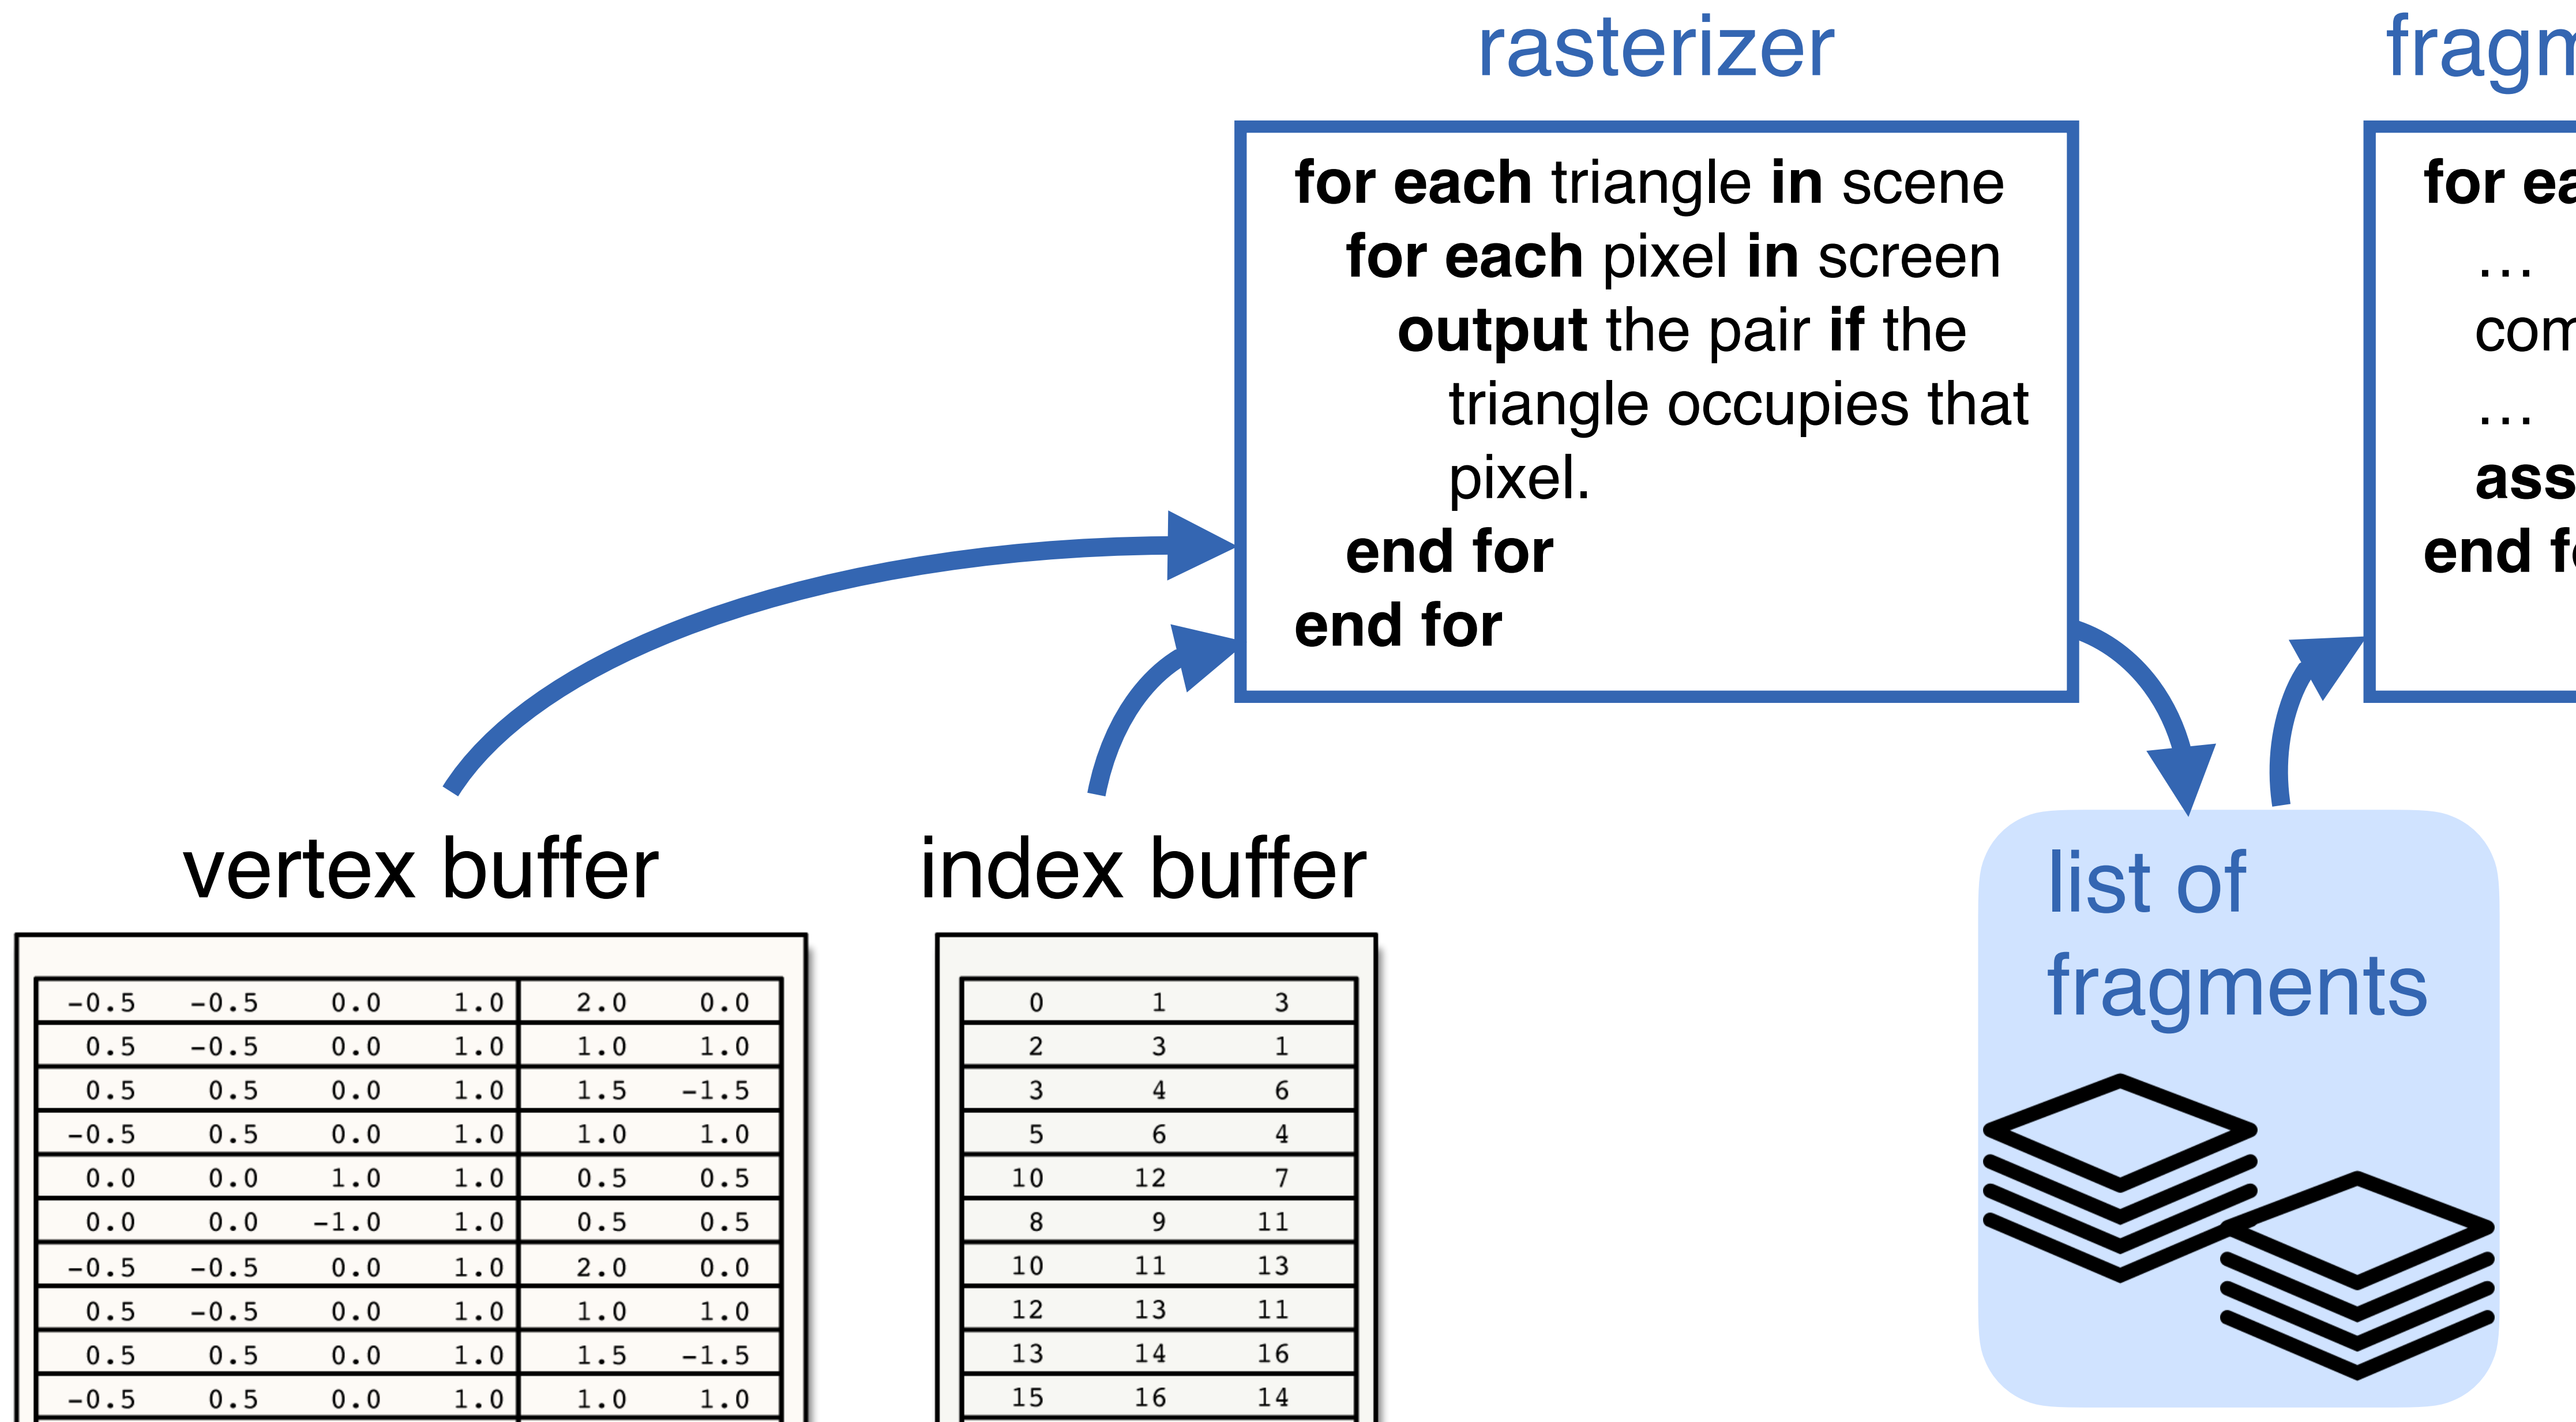
vertex

triangle	0	1	3
	2	3	1
	3	4	6
	5	6	4
	10	12	7
	8	9	11
	10	11	13
	12	13	11
	13	14	16
	15	16	14
	20	22	17
	18	19	21

Raster graphics pipeline




Raster graphics pipeline

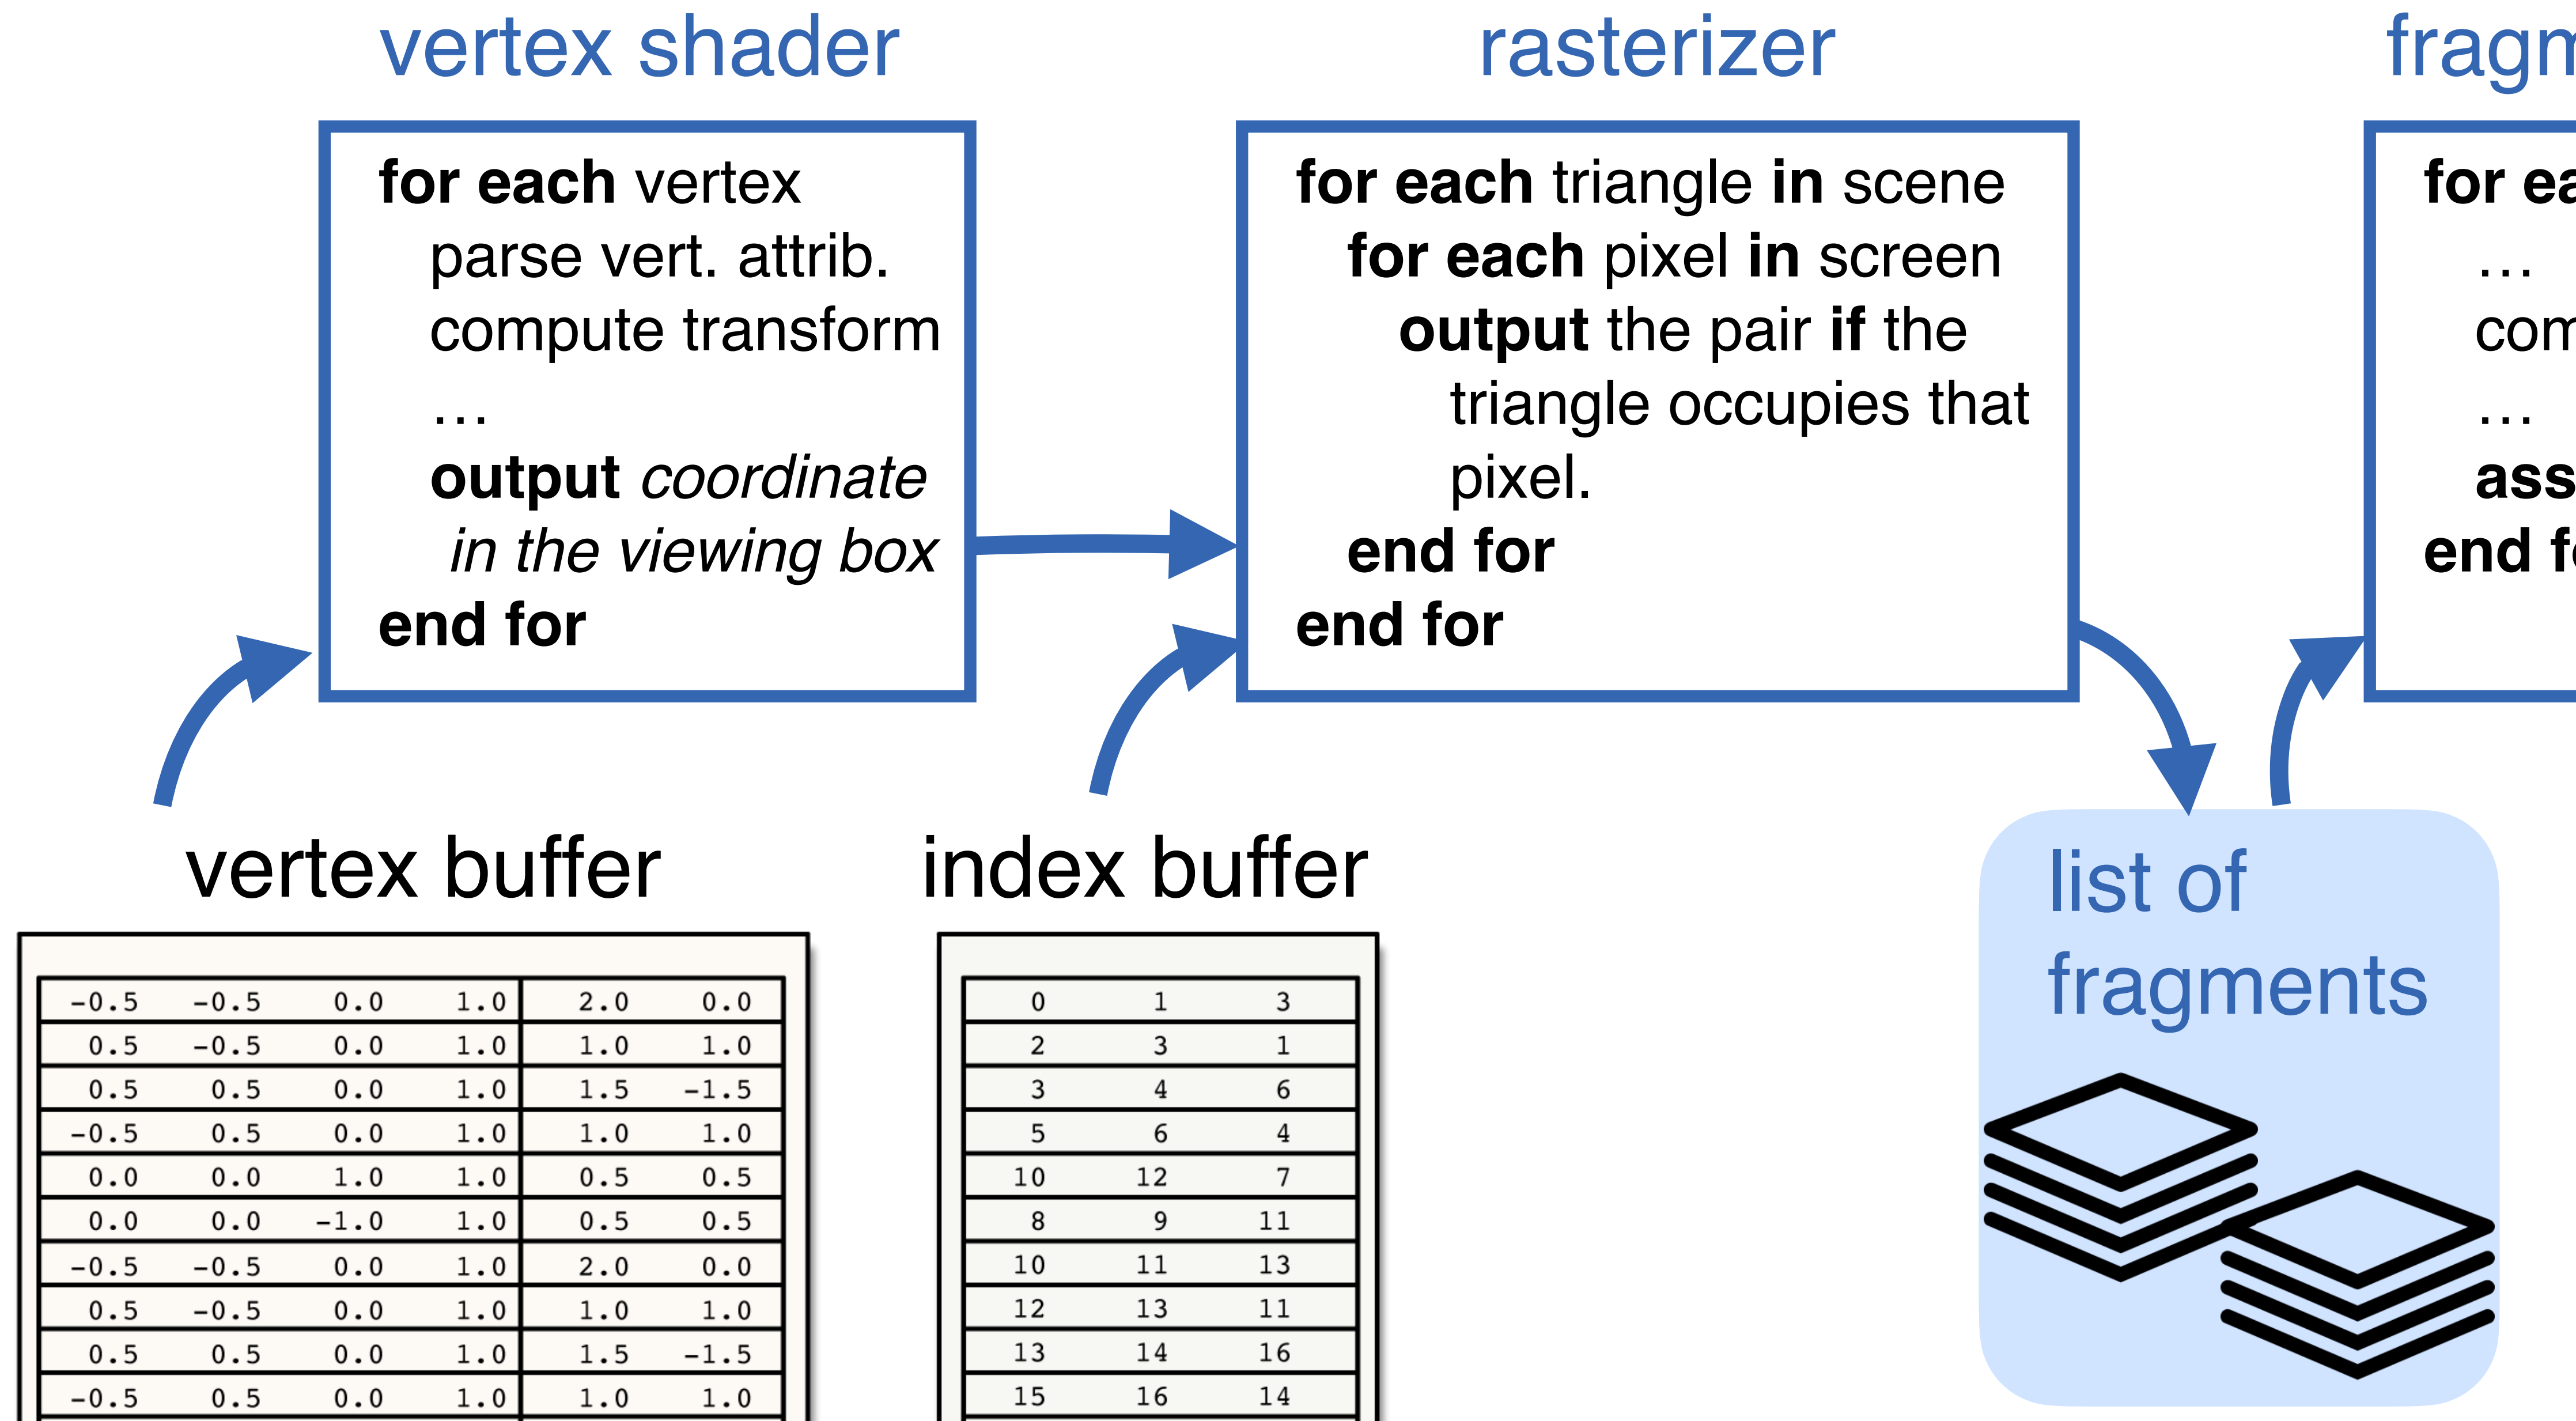


Raster graphics pipeline

Recall that before the rasterization, we need to perform some transformation...



Raster graphics pipeline



Raster graphics pipeline

vertex shader

```
for each vertex  
  parse vert. attrib.  
  compute transform  
  ...  
  output coordinate  
    in the viewing box  
end for
```

rasterizer

```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```


fragment shader

```
for each fragment  
  ...  
  compute color  
  ...  
  assign frag_color  
end for
```

resolve depth

```
for each pixel  
  keep only the  
  top fragment  
end for
```

color per
pixel



Raster graphics pipeline

vertex shader

```
for each vertex  
  parse vert. attrib.  
  compute transform  
  ...  
  output coordinate  
    in the viewing box  
end for
```

flexible

rasterizer

```
for each triangle in scene  
  for each pixel in screen  
    output the pair if the  
      triangle occupies that  
      pixel.  
  end for  
end for
```

fixed, hardcoded

fragment shader


```
for each fragment  
  ...  
  compute color  
  ...  
  assign frag_color  
end for
```

flexible

resolve depth

```
for each pixel  
  keep only the  
  top fragment  
end for
```

color per
pixel



- **Flexible:** Frequently reprogrammed depending on application
- **Fixed:** Hardcode in a chip
- Every **for each** loop can be parallelized

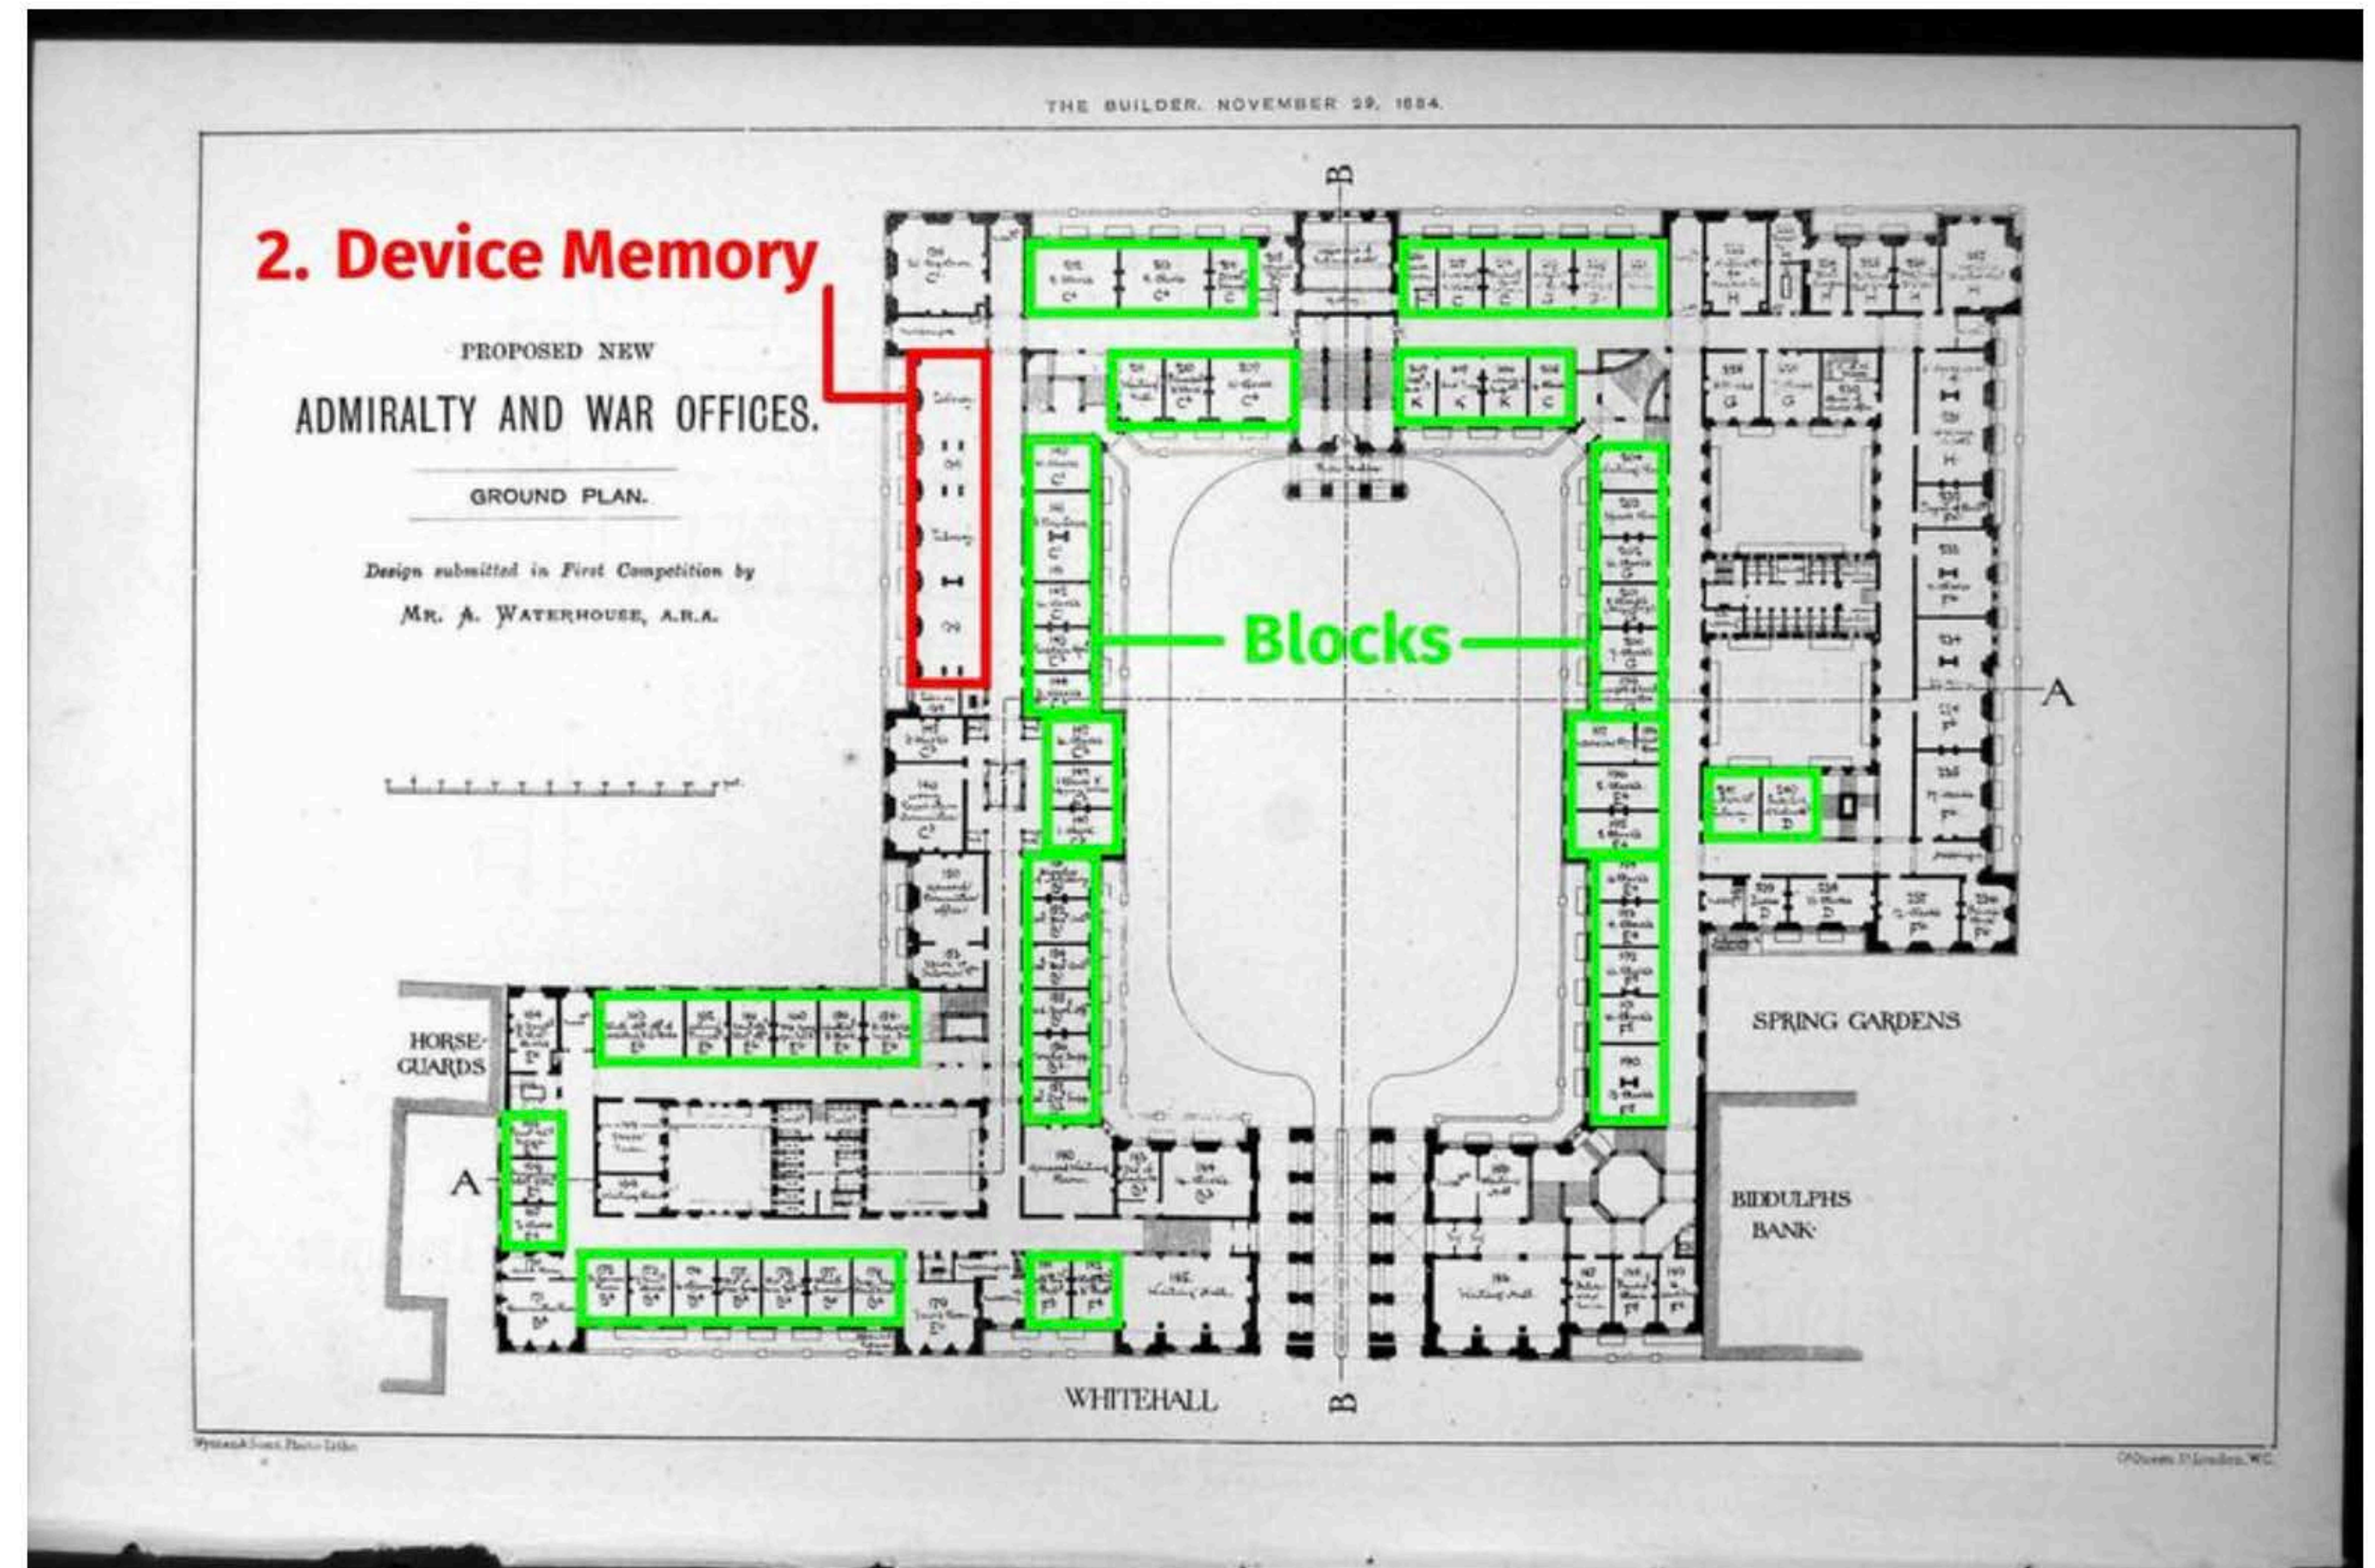
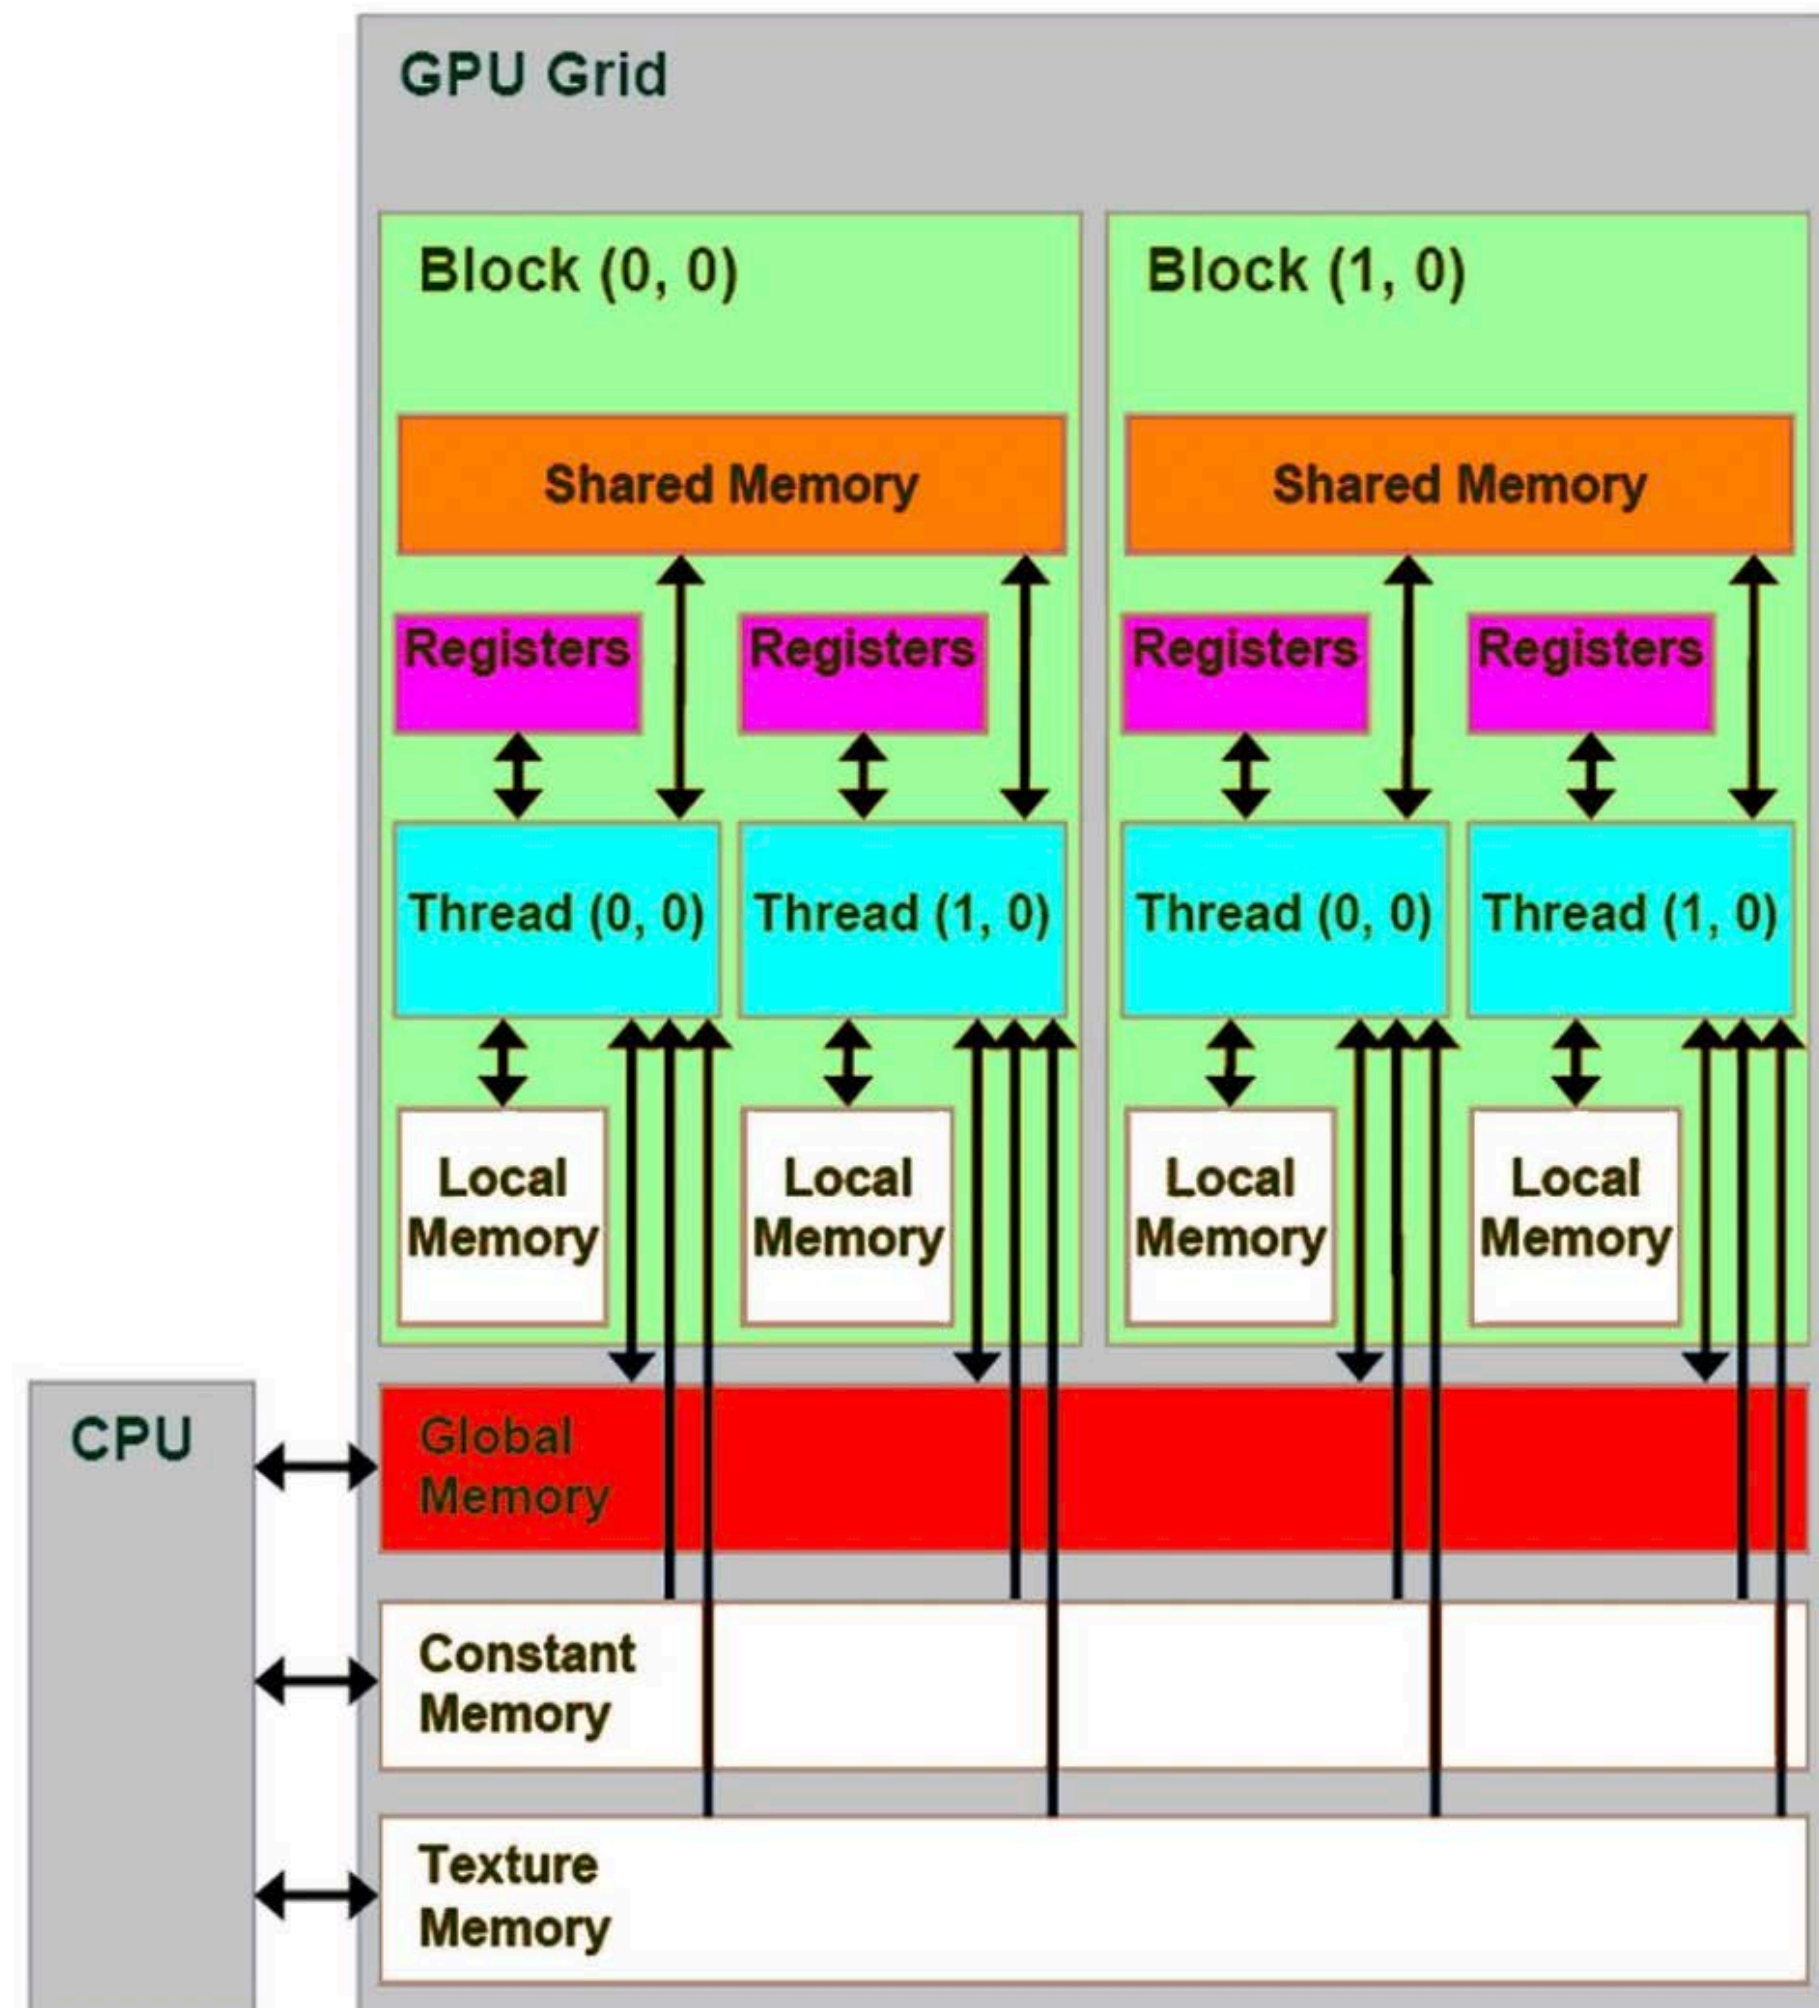
Graphics Processing Unit (GPU)

- Main algorithm
- Rasterization pipeline
- GPU
- Summary

What is GPU?

- Hardware that accelerates the graphics pipeline.
- GPU is a cluster of thousands of efficient workers knowing simple arithmetics working in parallel.
- GPU works in “Single Instruction Multiple Data” (SIMD) (as opposed to Single Instruction Single Data (SISD) like CPU)


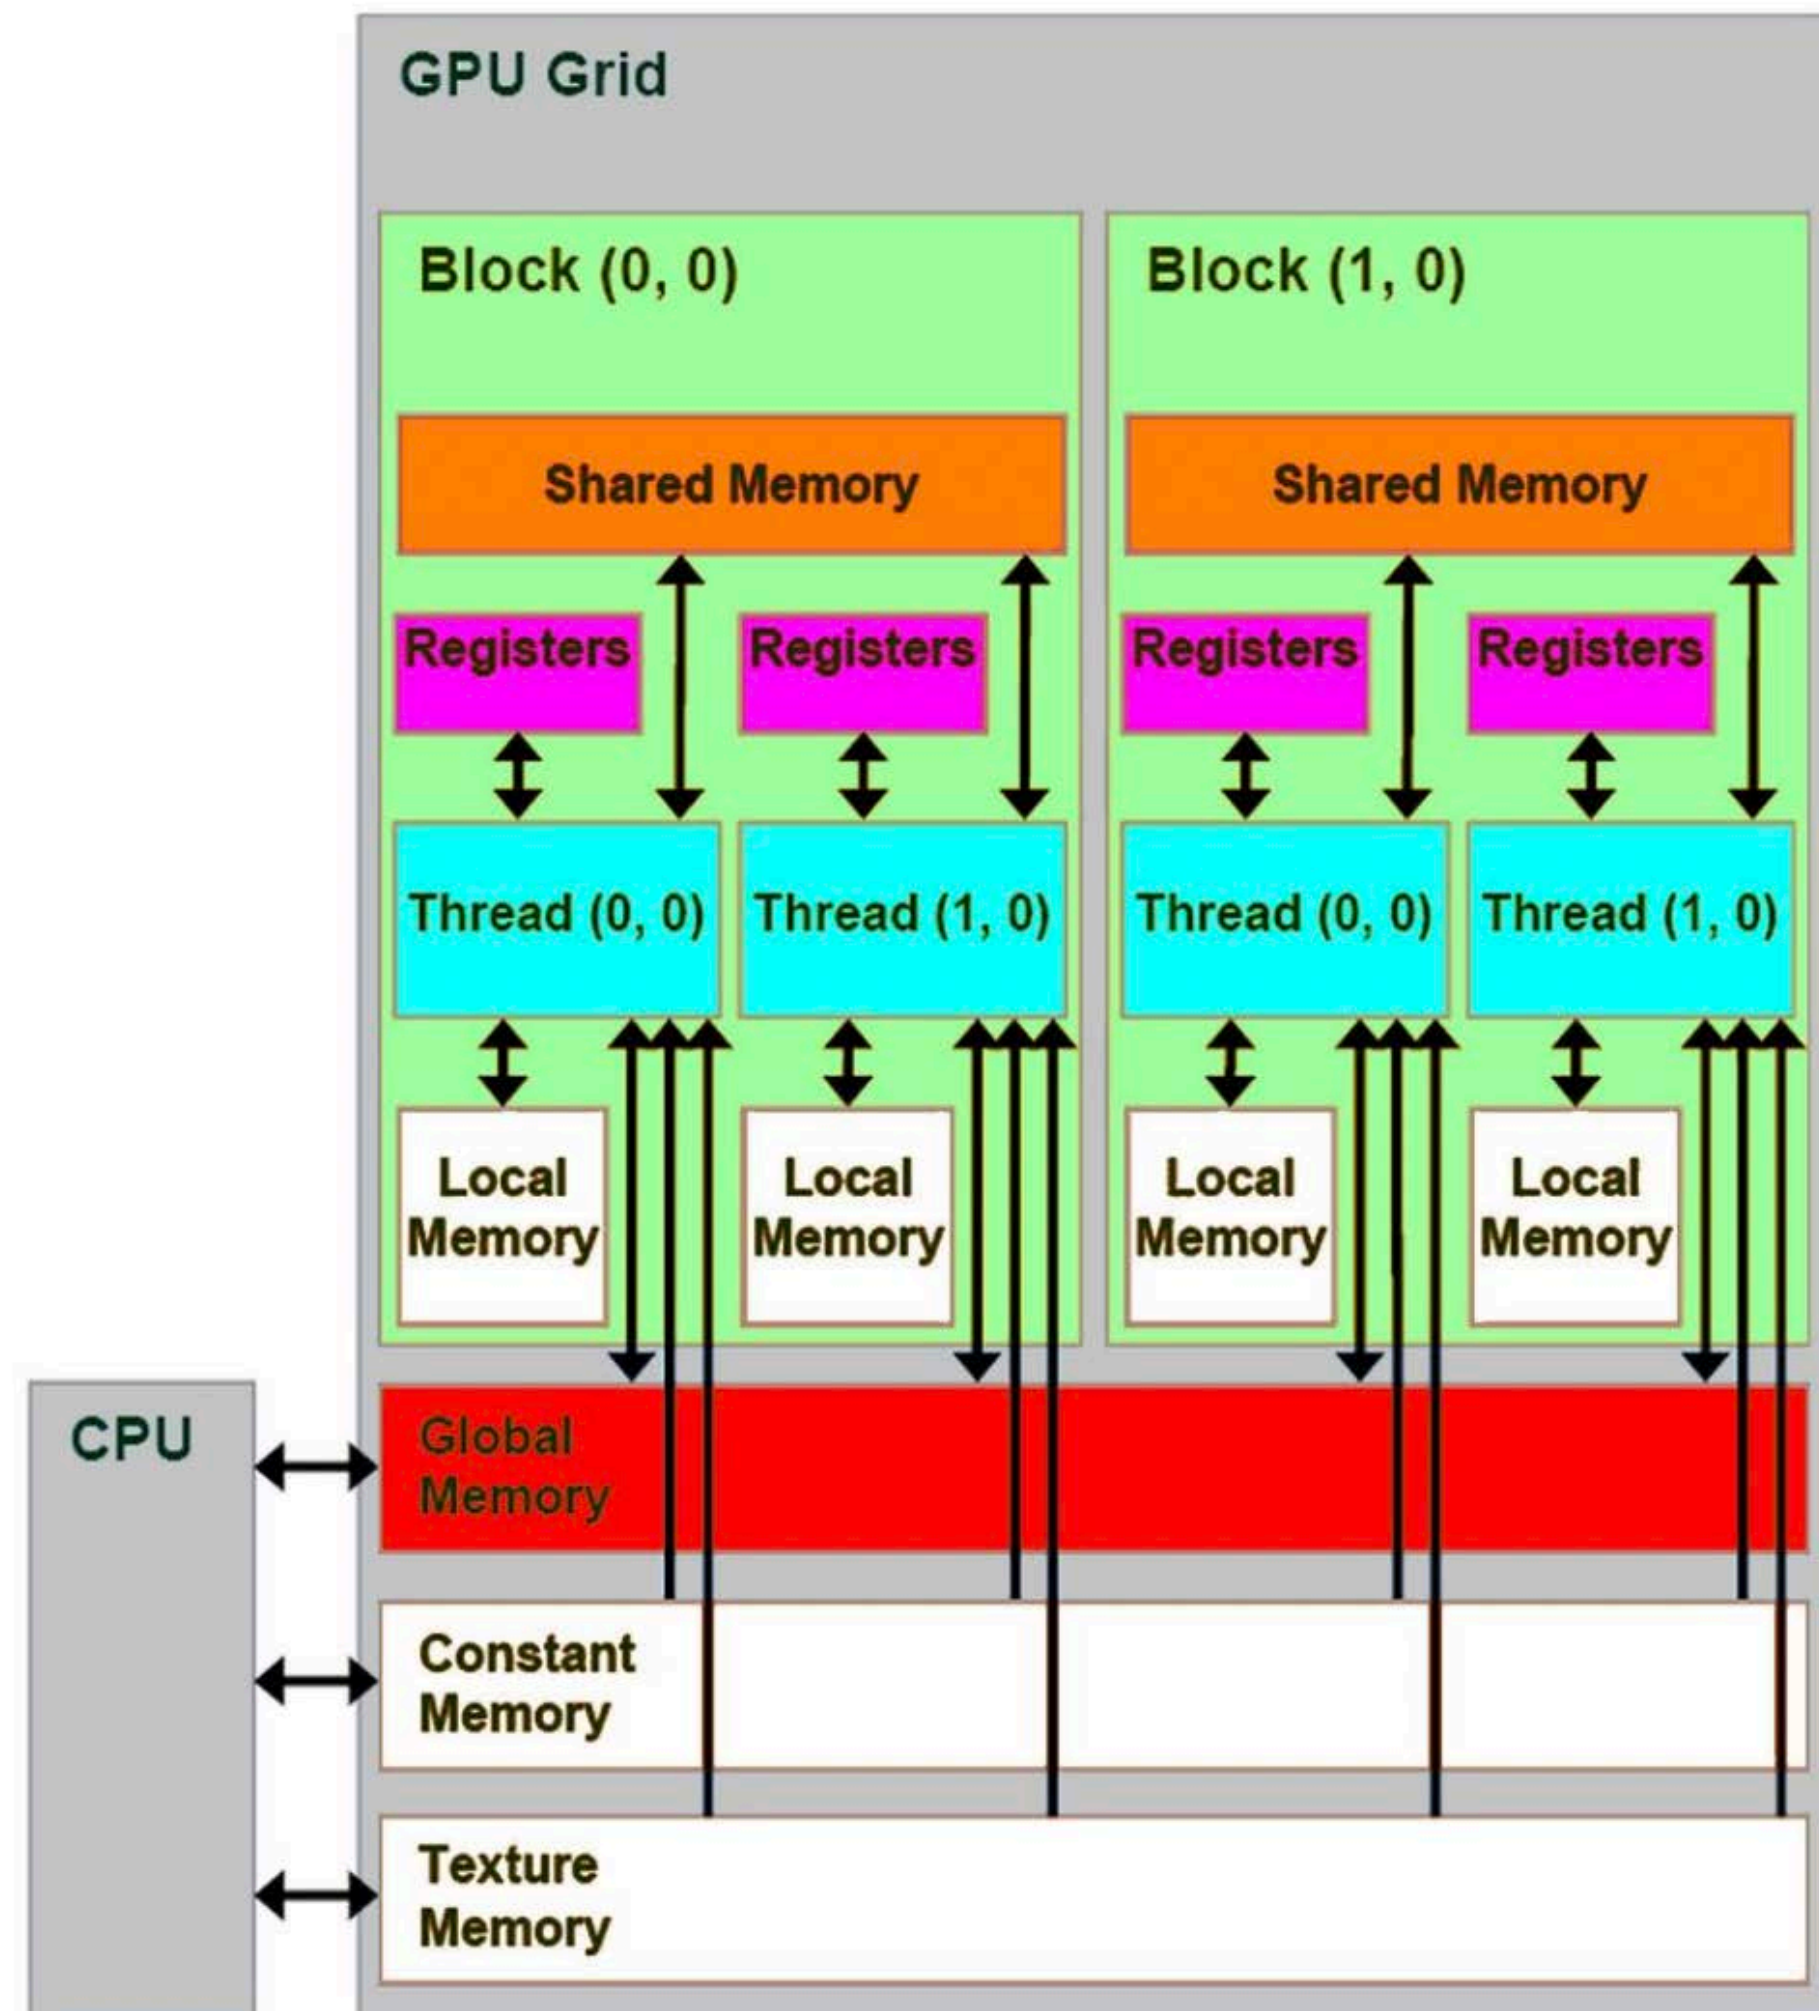
What is GPU?



A 19th century office building

(image courtesy Jean Feydy)

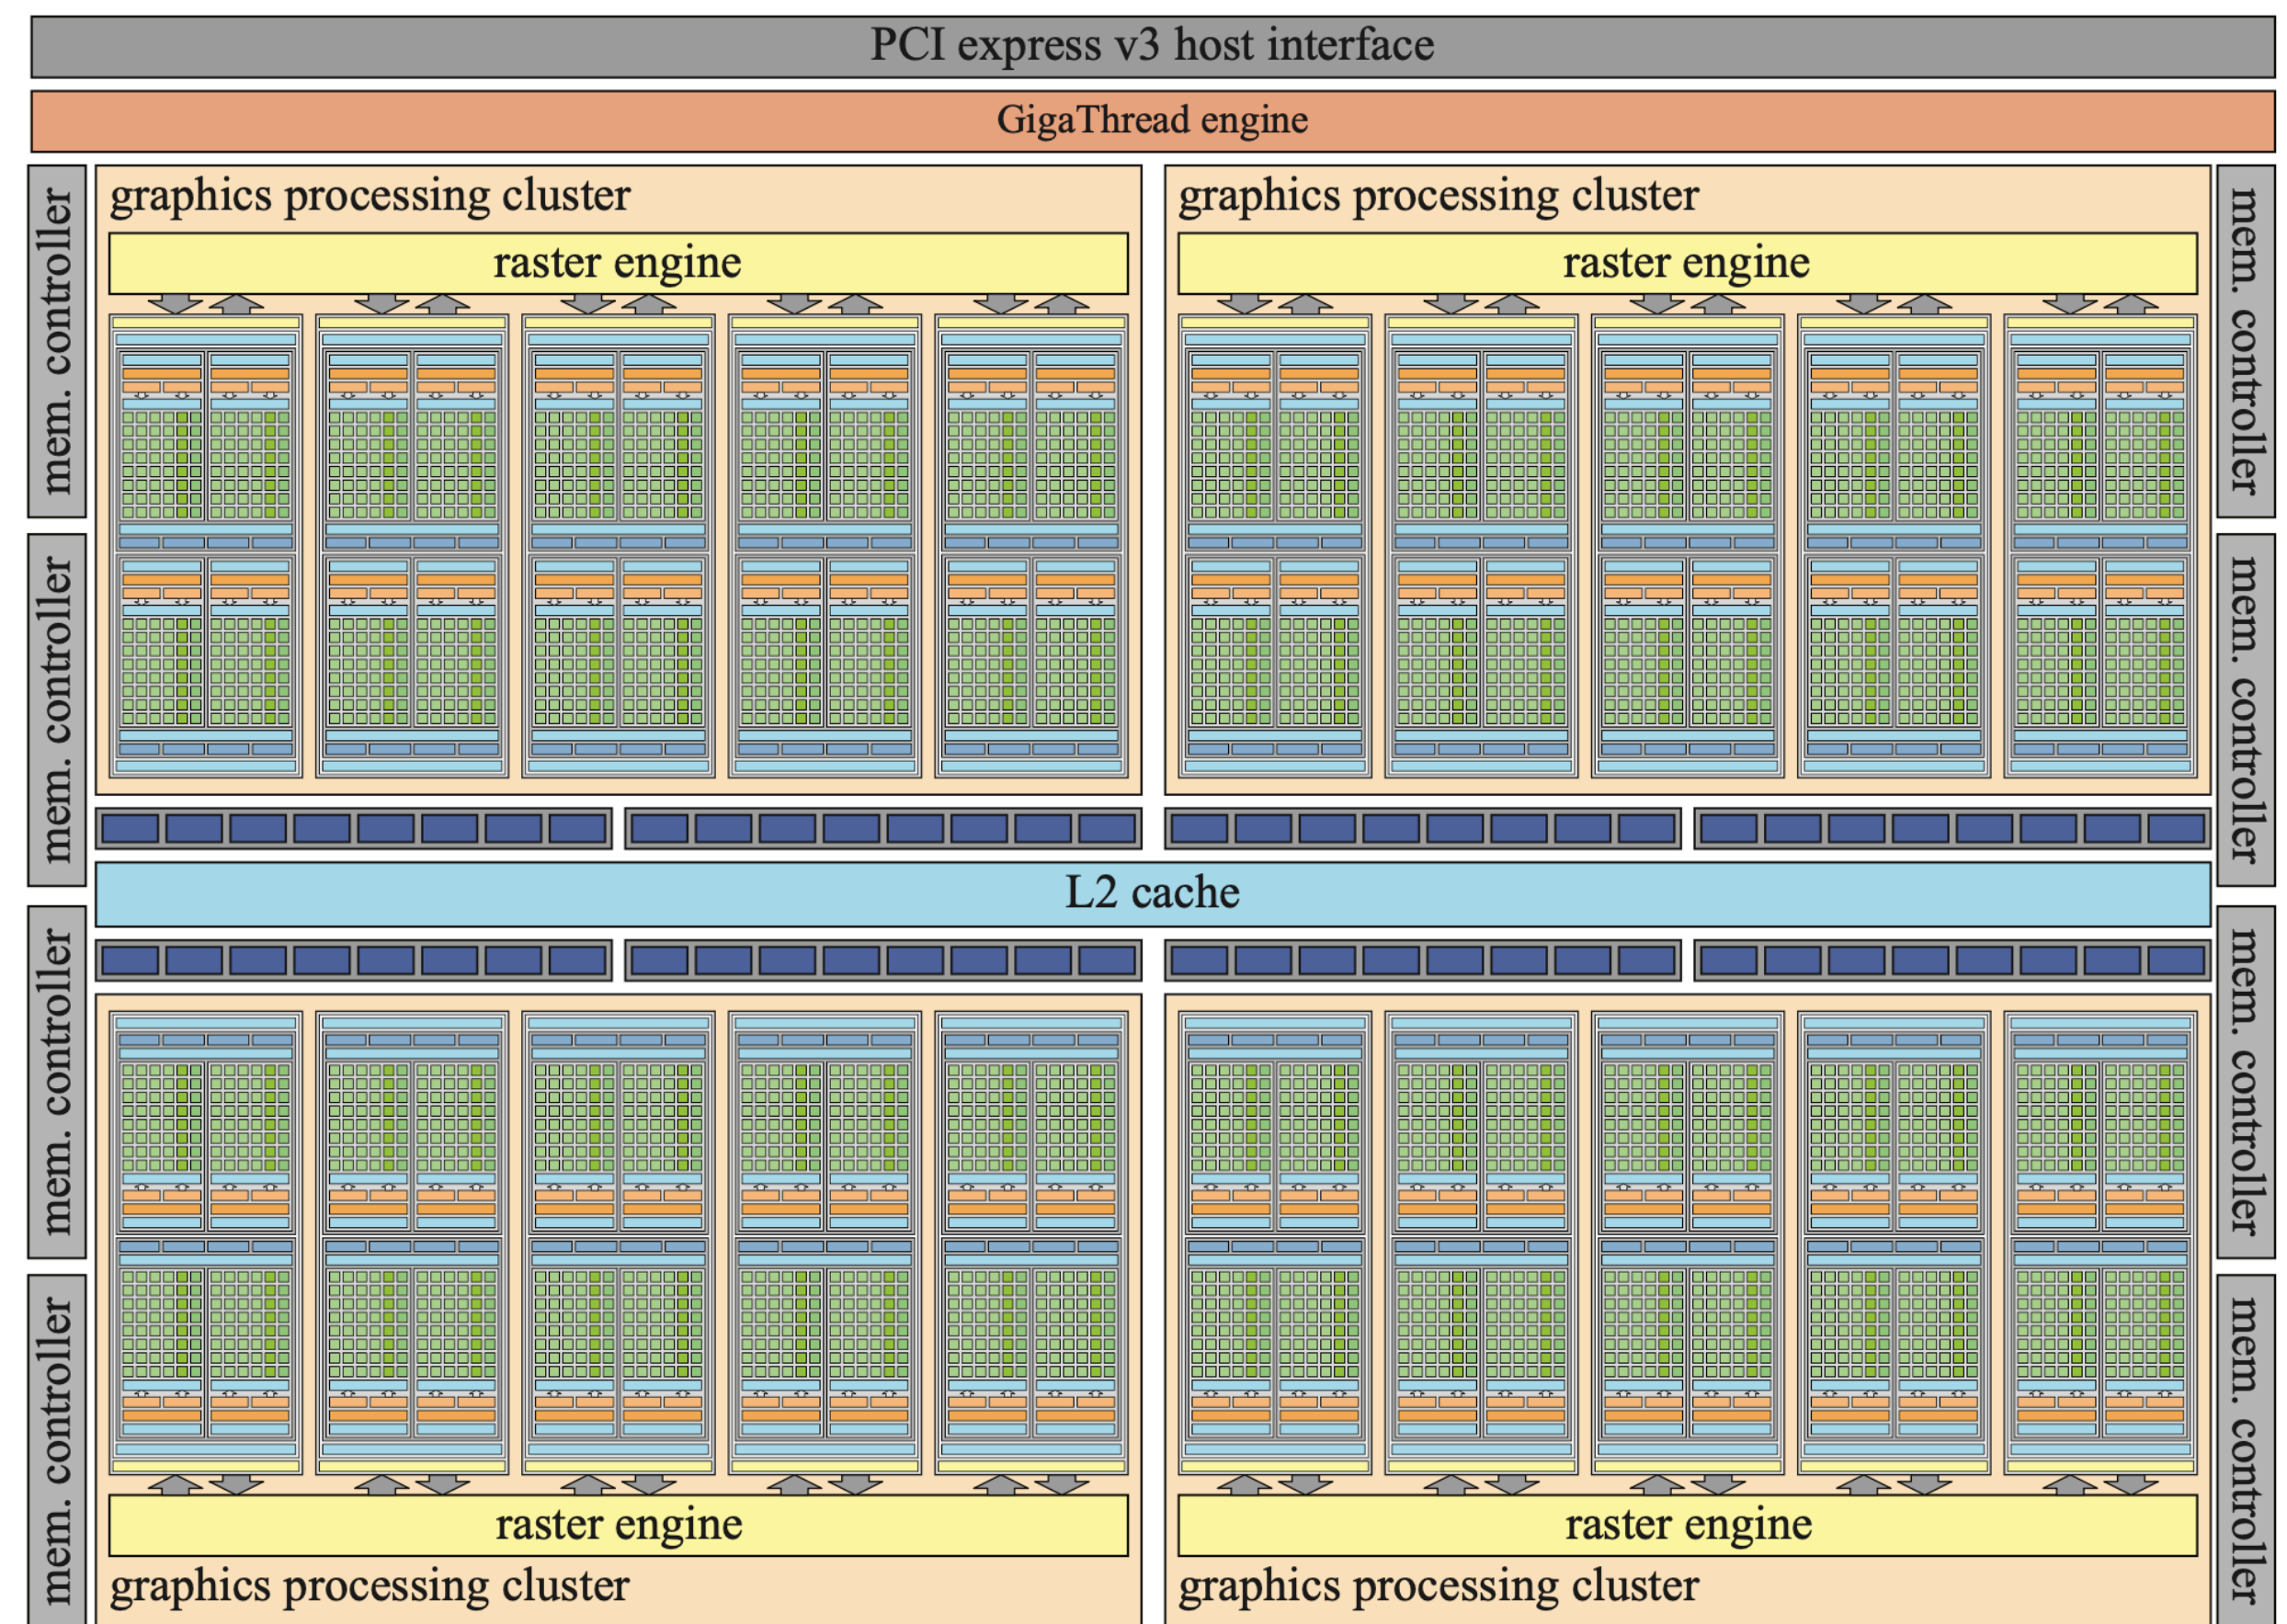
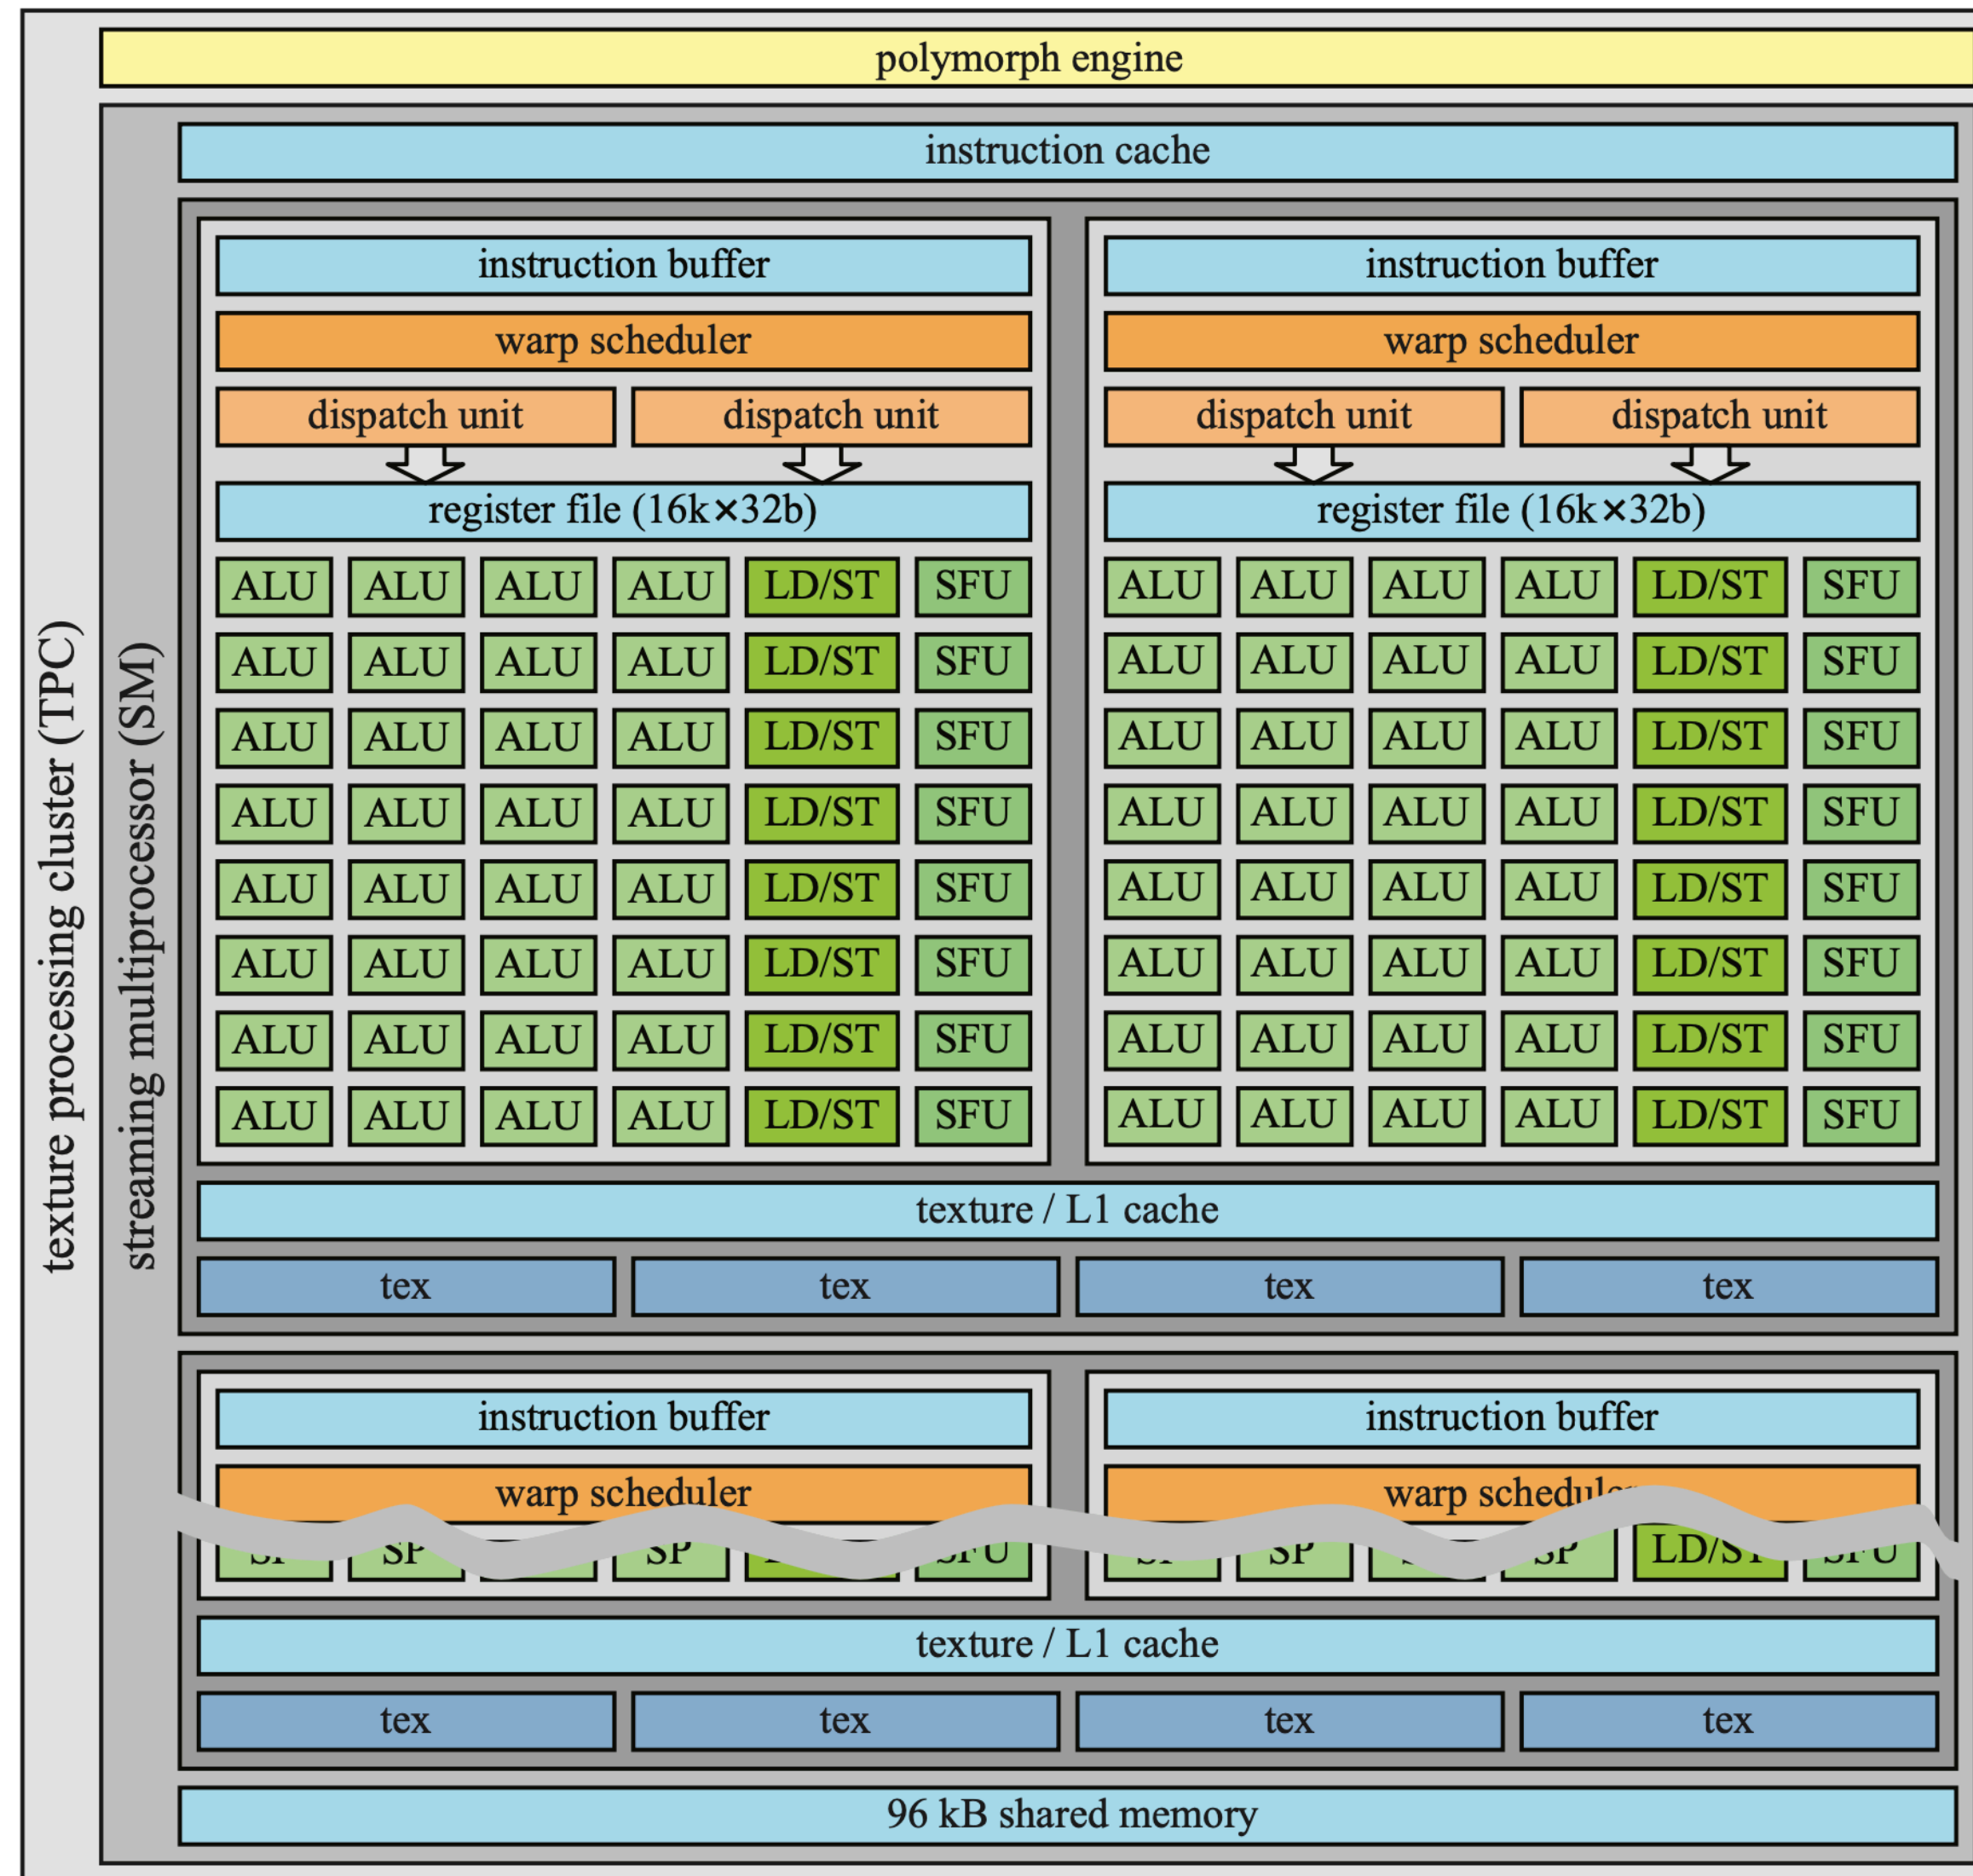
What is GPU?



A 19th century office building

(image courtesy Jean Feydy)

What is GPU?



(image from Real-Time Rendering Ch 23)

What is GPU?

- Frequent use of graphics pipeline \implies Dedicated hardware
- GPU manufacturers provide application programming interface (API) which is a list of functions allowing us to command GPU
- 90's: GPU are made with fixed functions (e.g. Legacy OpenGL)
- Mid 2000's: Programmable shaders (e.g. Modern OpenGL)

Definition A ***shader*** is a program that runs on GPU

- Trend: Towards programmability and flexibility
- General purpose GPU (e.g. CUDA), high performance computing, parallel computing, machine learning

Summary

- Main algorithm
- Rasterization pipeline
- GPU
- Summary

Summary

- **Fragments** are triangle-pixel incidence/intersection
- **Rasterization** and **ray tracing** are two ways of looping to find fragments
- A **buffer** is an allocated memory (on the graphics card)
- A **shader** is a program running on GPU
- Describe a shape by a geometry spreadsheet (vertex array object) stored in **vertex buffer** and **index buffer**
- Raster graphics pipeline:



- Implemented in GPU; **OpenGL** is an API to command it.