

## THE NON-IDENTITY PROBLEM

THERE is another question about personal identity. Each of us might never have existed. What would have made this true? The answer produces a problem that most of us overlook.

One of my aims in Part Four is to discuss this problem. My other aim is to discuss the part of our moral theory in which this problem arises. This is the part that covers how we affect future generations. This is the most important part of our moral theory, since the next few centuries will be the most important in human history.

### 119. HOW OUR IDENTITY IN FACT DEPENDS ON WHEN WE WERE CONCEIVED

What would have made it true that some particular person would never have existed? With one qualification, I believe

*The Time-Dependence Claim:* If any particular person had not been conceived when he was in fact conceived, it is *in fact* true that he would never have existed.

This claim is not obviously true. Thus one woman writes:

It is always fascinating to speculate on who we would have been if our parents had married other people.<sup>1</sup>

In wondering who she would have been, this woman ignores the answer: 'No one'.

Though the Time-Dependence Claim is not obviously true, it is not controversial, and it is easy to believe. It is thus unlike the Reductionist View about personal identity over time. This is one of several competing views, and is hard to believe. The Time-Dependence Claim is not about personal identity over time. It is about a different though related subject: personal identity in different possible histories of the world. Several views about this subject are worth discussing. But the Time-Dependence Claim is *not* one of these views. It is a claim that is true on *all* of these views.

As I have said, the claim should be qualified. Each of us grew from a particular pair of cells. Each pair is an ovum and the spermatozoon which, out of millions, fertilized this ovum. Suppose that my mother had not

conceived a child at the time when in fact she conceived me. And suppose that she had conceived a child within a few days of this time. This child would have grown from the same particular ovum from which I grew. But even if this child had been conceived only a few seconds earlier or later, it is almost certain that he would have grown from a different spermatozoon. This child would have had some but not all of my genes. Would this child have been me?

We are inclined to believe that any question about our identity must have an answer, which must be either Yes or No. As before, I reject this view. There are cases in which our identity is indeterminate. What I have just described may be such a case. If it is, my question has no answer. It is neither true nor false that, if these events had occurred, I would never have existed. Though I can always ask, 'Would I have existed?', this would here be an empty question.

These last claims are controversial. Since I want my Time-Dependence Claim not to be controversial, I shall set aside these cases. The claim can become

(TD2) If any particular person had not been conceived within a month of the time when he was in fact conceived, he would in fact never have existed.

I claim that this is *in fact* true. I do *not* claim that it is *necessarily* true. The different views about this subject make competing claims about what is necessary. It is because I claim less that my claim is not controversial. Those who disagree about what *could* have happened may agree about what *would in fact* have happened. As I shall argue, the holders of all plausible views would agree with me.

These views make claims about the *necessary properties* of each particular person. Some of a person's necessary properties are had by everyone: these are the properties that are necessary to being a person. What concerns us here are the *distinctive* necessary properties of each particular person. Suppose I claim that *P* is one of Kant's distinctive necessary properties. This means that Kant could not have lacked *P*, and that only Kant could have had *P*.

According to

*The Origin View*, each person has this distinctive necessary property: that of having grown from the particular pair of cells from which this person in fact grew.<sup>2</sup>

This property cannot be *fully* distinctive. Any pair of identical twins *both* grew from such a pair of cells. And any fertilized ovum might have later split, and produced twins. The Origin View must be revised to meet this problem. But I need not discuss this revision. It is enough for my purposes that, on this view, Kant could not have grown from a different pair of cells.

It is irrelevant that, because there can be twins, it is false that *only* Kant could have grown from this pair of cells.

Holders of the Origin View would accept my claim that, if Kant had not been conceived within a month of the time when he was conceived, he would in fact never have existed. If he had not been conceived in that month, no child would in fact have grown from the particular pair of cells from which he grew. (This claim makes an assumption both about the distinctive necessary properties of this pair of cells, and about the human reproductive system. But these assumptions are not controversial.)

According to certain other views, Kant could have grown from a different pair of cells. On

*The Featureless Cartesian View*, Kant was a particular Cartesian Ego, which had *no* distinctive necessary properties.

On this view, a person's identity has no connections with his physical and mental characteristics. Kant might have been me, and vice versa, though, if this had happened, no one would have noticed any difference. It is at worst mildly controversial to claim, as I did, that we should reject *this* version of the Cartesian View.

Two other views are closely related. On

*The Descriptive View*, each person has several distinctive necessary properties. These are this person's most important distinctive properties, and they do not include having grown from a particular pair of cells.

In the case of Kant, these properties would include his authorship of certain books. One version of this view does not claim that Kant must have had *all* these properties. Anyone with most of these properties would have been Kant.

On

*The Descriptive Name View*, every person's name means 'the person who . . .'. For us now, 'Kant' means 'the person who wrote the *Critique of Pure Reason*, etc'. A particular person's necessary properties are those that would be listed when we explain the meaning of this person's name.

Both this and the Descriptive View might be combined with the other version of Cartesianism. Kant might be claimed to be the Cartesian Ego whose distinctive necessary properties include the authorship of certain books. But the two Descriptive Views need not add this claim.<sup>3</sup>

One objection to the Descriptive Views is that each person's life could have been very different. Kant could have died in his cradle. Since this is possible, the authorship of certain books cannot be one of Kant's necessary properties.

One reply to this objection retreats to a weaker claim. It could be said:

Though this property is not necessary, it is distinctive. Kant might not have written these books. But, in any possible history in which a single person wrote these books, this person would have been Kant.

I need not discuss whether this, or some other reply, meets this objection. Even if the objection can be met, my Time-Dependence Claim is true.

On both Descriptive Views, Kant could have grown from a different pair of cells, or even had different parents. This would have happened if Kant's mother had not conceived a child when she conceived him, and some other couple had conceived a child who later wrote the *Critique of Pure Reason*, etc. On the Descriptive Views, this child would have been Kant. He would not have been *called* Kant. But this does not worry holders of these views. They would claim that, if this had happened, Kant would have had both different parents and a different name.

Though they believe that this might have happened, most holders of the Descriptive Views would accept my claim that it would *not in fact* have happened. If they claim that it *would* have happened, they must accept an extreme version of Tolstoy's view, stated in the epilogue of *War and Peace*, that history does not depend on the decisions made by particular people. On this view, if Napoleon's mother had remained childless, history would have provided a 'substitute Napoleon', who would have invaded Russia in 1812. And, if Kant's mother had remained childless, history would have provided another author of the *Critique of Pure Reason*. This view is too implausible to be worth discussing.

There is another way in which holders of the Descriptive Views might reject my claim. They might claim that Kant's necessary properties were far less distinctive. They might for instance merely be: being his mother's first child. This claim meets the objection that each person's life might have been very different. But this claim is also too implausible to be worth discussing. I am the second of my mother's three children. This claim implies absurdly that, if my mother had conceived no child when she in fact conceived me, I would have been my younger sister.

Consider next the possible history in which the Descriptive Views seem most plausible. Suppose that Kant's mother had not conceived a child when she conceived him, and that one month later she conceived a child who was exactly like Kant. This child would have grown from a different pair of cells; but by an amazing coincidence, of a kind that never actually happens, this child would have had all of Kant's genes. And suppose that, apart from the fact and the effects of being born later, this child would have lived a life that was just like Kant's, writing the *Critique of Pure Reason*, etc.

On the Descriptive Views, this child would have been Kant. Holders of the Origin View might object:

Kant was a particular person. In your imagined possible history, you have not shown that you are referring to *this* particular person. In this imagined history, there would have been someone who was *exactly like*

this person. But exact similarity is not the same as numerical identity, as is shown by any two exactly similar things.

These remarks explain why the Origin View refers to the particular pair of cells from which a person grew.

A fifth view also makes such a direct reference. On

*The Backward Variation View*, this reference need not be to the point of origin, or to the cells from which a person grew. The reference can be made to any person at any time in his life. By making such a reference, we can describe how this person might have had a different origin.<sup>4</sup>

Consider a holder of this view who, in 1780, is attending one of Kant's lectures. This person might claim:

Kant is the person standing *there*, looking just like *that*. Kant might have had different parents, and lived a different life up until the recent past. For this to have been what happened, all that is needed is that this different life would have led Kant to be now standing *there*, looking just like *that*.

This view must make some further claims. But it meets the objection that, to justify a claim of identity, we need more than similarity. Holders of the Origin View therefore need a different objection to the Backward Variation View. For my purposes, I need not decide between these views.

On the Backward Variation View, Kant might have had a different origin. But holders of this view would accept my claim that, in fact, this would not have happened. They would agree that, if Kant had not been conceived within a month of the time when he was conceived, he would in fact never have existed.

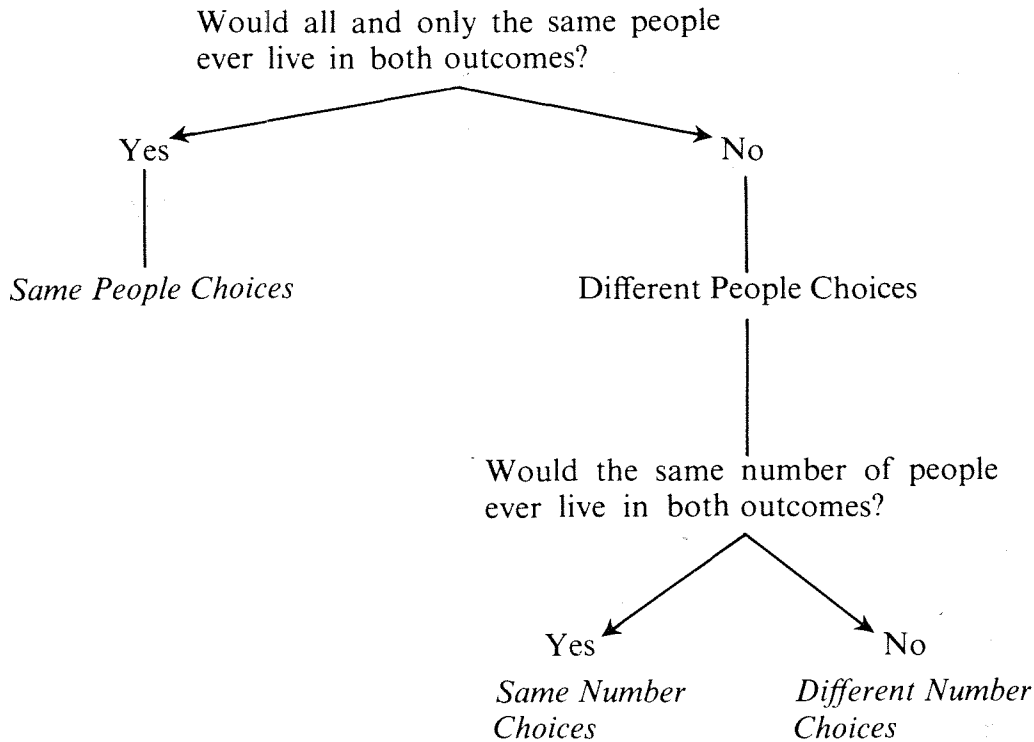
I have now described all of the views about our identity in different possible histories.<sup>5</sup> I discuss in endnote 6 how these views are related to the different views about our identity over time. On all of the plausible views, my Time-Dependence Claim is true. This claim applies to everyone. You were conceived at a certain time. It is in fact true that, if you had not been conceived within a month of that time, *you* would never have existed.

## 120. THE THREE KINDS OF CHOICE

Unless we, or some global disaster, destroy the human race, there will be people living later who do not now exist. These are *future people*. Science has given to our generation great ability both to affect these people, and to predict these effects.

Two kinds of effect raise puzzling questions. We can affect the identities of future people, or *who* the people are who will later live. And we can affect the number of future people. These effects give us different kinds of choice.

In comparing any two acts, we can ask:



Different Number Choices affect both the number and the identities of future people. Same Number Choices affect the identities of future people, but do not affect their number. Same People Choices affect neither.

#### 121. WHAT WEIGHT SHOULD WE GIVE TO THE INTERESTS OF FUTURE PEOPLE?

Most of our moral thinking is about Same People Choices. As I shall argue, such choices are not as numerous as most of us assume. Very many of our choices will in fact have some effect on both the identities and the number of future people. But in most of these cases, because we cannot predict what the particular effects would be, these effects can be morally ignored. We can treat these cases as if they were Same People Choices.

In some cases we can predict that some act either may or will be against the interests of future people. This can be true when we are making a Same People Choice. In such a case, whatever we choose, all and only the same people will ever live. Some of these people will be future people. Since these people will exist whatever we choose, we can either harm or benefit these people in a quite straightforward way.

Suppose that I leave some broken glass in the undergrowth of a wood. A hundred years later this glass wounds a child. My act harms this child. If I had safely buried the glass, this child would have walked through the wood unharmed.

Does it make a moral difference that the child whom I harm does not now exist?

On one view, moral principles cover only people who can *reciprocate*, or harm and benefit each other. If I cannot be harmed or benefited by this child, as we can plausibly suppose, the harm that I cause this child has no moral importance. I assume that we should reject this view.<sup>7</sup>

Some writers claim that, while we ought to be concerned about effects on future people, we are morally justified in being less concerned about effects in the further future. This is a common view in welfare economics, and cost-benefit analysis. On this view, we can *discount* the more remote effects of our acts and policies, at some rate of  $n$  per cent per year. This is called the *Social Discount Rate*.

Suppose we are considering how to dispose safely of the radio-active matter called *nuclear waste*. If we believe in the Social Discount Rate, we shall be concerned with safety only in the nearer future. We shall not be troubled by the fact that some nuclear waste will be radio-active for thousands of years. At a discount rate of five per cent, one death next year counts for more than a billion deaths in 500 years. On this view, catastrophes in the further future can now be regarded as morally trivial.

As this case suggests, the Social Discount Rate is indefensible. Remoteness in time roughly correlates with some important facts, such as predictability. But, as I argue in Appendix F, these correlations are too rough to justify the Social Discount Rate. The present moral importance of future events does *not* decline at a rate of  $n$  per cent per year. Remoteness in time has, in itself, no more significance than remoteness in space. Suppose that I shoot some arrow into a distant wood, where it wounds some person. If I should have known that there might be someone in this wood, I am guilty of gross negligence. Because this person is far away, I cannot identify the person whom I harm. But this is no excuse. Nor is it any excuse that this person is far away. We should make the same claims about effects on people who are temporally remote.

## 122. A YOUNG GIRL'S CHILD

Future people are, in one respect, unlike distant people. We can affect their identity. And many of our acts have this effect.

This fact produces a problem. Before I describe this problem, I shall repeat some preliminary remarks. I assume that one person can be worse off than another, in morally significant ways, and by more or less. But I do not assume that these comparisons could be, even in principle, precise. I assume that there is only rough or partial comparability. On this assumption, it could be true of two people that neither is worse off than the other, but this would not imply that these people are exactly equally well off.

'Worse off' could be taken to refer, either to someone's level of happiness, or more narrowly to his standard of living, or, more broadly, to

the quality of his life. Since it is the broadest, I shall often use the phrase 'the quality of life'. I also call certain lives 'worth living'. This description can be ignored by those who believe that there could not be lives that are not worth living. But, like many other people, I believe that there are such lives. Finally, I extend the ordinary use of the phrase 'worth living'. If one of two people would have a lower quality of life, I call his life to this extent 'less worth living'.

When considering future people, we must answer two questions:

- (1) If we cause someone to exist, who will have a life worth living, do we thereby benefit this person?
- (2) Do we also benefit this person if some act of ours is a remote but necessary part of the cause of his existence?

These are difficult questions. If we answer Yes to both, I shall say that we believe *that causing to exist can benefit*.

Some people answer Yes to (1) but No to (2). These people give their second answer because they use 'benefit' in its ordinary sense. As I argued in Section 25, we ought for moral purposes to extend our use of 'benefit'. If we answer Yes to (1) we should answer Yes to (2).

Many people answer No to both these questions. These people might say: 'We benefit someone if it is true that, if we had not done what we did, this would have been worse for this person. If we had not caused someone to exist, this would *not* have been worse for this person.'

I believe that, while it is defensible to answer No to both these questions, it is also defensible to answer Yes to both. For those who doubt this second belief I have written Appendix G. Since I believe that it is defensible both to claim and to deny that causing to exist can benefit, I shall discuss the implications of both views.

Consider

*The 14-Year-Old Girl.* This girl chooses to have a child. Because she is so young, she gives her child a bad start in life. Though this will have bad effects throughout this child's life, his life will, predictably, be worth living. If this girl had waited for several years, she would have had a different child, to whom she would have given a better start in life.

Since such cases are, at least in the United States, becoming very numerous, they raise a practical problem.<sup>8</sup> They also raise a theoretical problem.

Suppose that we tried to persuade this girl that she ought to wait. We claimed: 'If you have a child now, you will soon regret this. If you wait, this will be better for you.' She replied: 'This is my affair. Even if am doing what will be worse for me, I have a right to do what I want.'

We replied: 'This is not entirely your affair. You should think not only of



yourself, but also of your child. It will be worse for him if you have him now. If you have him later, you will give him a better start in life.'

We failed to persuade this girl. She had a child when she was 14, and, as we predicted, she gave him a bad start in life. Were we right to claim that her decision was worse for her child? If she had waited, this particular child would never have existed. And, despite its bad start, his life is worth living. Suppose first that we do *not* believe that causing to exist can benefit. We should ask, 'If someone lives a life that is worth living, is this worse for this person than if he had never existed?' Our answer must be No. Suppose next that we believe that causing to exist *can* benefit. On this view, this girl's decision benefits her child.

On both views, this girl's decision was not worse for her child. When we see this, do we change our mind about this decision? Do we cease to believe that it would have been better if this girl had waited, so that she could give to her first child a better start in life? I continue to have this belief, as do most of those who consider this case. But we cannot defend this belief in the natural way that I suggested. We cannot claim that this girl's decision was worse for her child. What is the objection to her decision? This question arises because, in the different outcomes, different people would be born. I shall therefore call this the *Non-Identity Problem*.<sup>9</sup>

It may be said:

In one sense, this girl's decision *was* worse for her child. In trying to persuade this girl not to have a child now, we can use the phrase 'her child' and the pronoun 'he' to cover *any* child that she might have. These words need not refer to one particular child. We can truly claim: 'If this girl does not have her child now, but waits and has him later, *he* will not be the same particular child. If she has him later, he will be a different child.' By using these words in this way, we can explain why it would be better if this girl waits. We can claim:

- (A) The objection to this girl's decision is that it will probably be worse for her child. If she waits, she would give him a better start in life.

Though we can truly make this claim, it does *not* explain the objection to this girl's decision. This becomes clear after the girl has had her child. The phrase 'her child' now naturally refers to this particular child. And this girl's decision was *not* worse for *this* child. Though there is a sense in which (A) is true, (A) does not appeal to a familiar moral principle.

On one of our familiar principles, it is an objection to someone's choice that this choice will be worse for, or be against the interests of, any other particular person. If we claim that this girl's decision was worse for her child, we cannot be claiming that it was worse for a particular person. We cannot claim, of the girl's child, that her decision was worse for *him*. We must admit that, in claim (A), the words 'her child' do not refer to her child.

(A) is not about what is good or bad for any of the particular people who ever live. (A) appeals to a new principle, that must be explained and justified.

If (A) seems to appeal to a familiar principle, this is because it has two senses. Here is another example. A general shows military skill if, in many battles, he always makes his the winning side. But there are two ways of doing this. He might win victories. Or he might always, when he is about to lose, change sides. A general shows no military skill if it is only in the second sense that he always makes his the winning side.

To what principle does (A) appeal? We might suggest

(B) It is bad if people are worse off than they would have been.

But this again has two senses, since 'they' could be taken to refer to the same particular people. We should remove this ambiguity. And we should state the principle in a way that shows the kind of choice to which it applies. These are Same Number Choices, which affect the identities of future people, but do not affect their number. We might suggest

*The Same Number Quality Claim, or Q:* If in either of two outcomes the same number of people would ever live, it would be bad if those who live are worse off, or have a lower quality of life, than those who would have lived.

This claim is plausible. And it implies what we believe about the 14-Year-Old Girl. The child that she has now will probably be worse off than a child she could have had later would have been, since this other child would have had a better start in life. If this is true, Q implies that this girl's choice has a bad effect. Q implies that it would have been better if she had waited, and had a child later.

We may shrink from claiming, of this girl's actual child, that it would have been better if he had never existed. But, if we claimed earlier that it would be better if this girl waits, this is what we must claim. We cannot consistently make a claim and deny this same claim later. If (1) in 1990 it *would be* better if this girl waits and has a child later, then (2) in 2020 it *would have been* better if she had waited and had a child later. And (2) implies (3) that it would have been better if the child who existed had not been her actual child. If we cannot accept (3), we must reject (1).

I suggest that, on reflection, we can accept (3). I believe that, if I was the actual child of this girl, I could accept (3). (3) does not imply that my existence is *bad*, or intrinsically morally undesirable. The claim is merely that, since a child born later would probably have had a better life than mine, it would have been better if my mother had waited, and had a child later. This claim need not imply that I ought rationally to regret that my mother had *me*, or that she ought rationally to regret this. Since it would have been better if she had waited, she ought perhaps to have some moral regret. And it is probably true that she made the outcome worse for herself.

But, even if this is true, it does not show that she ought rationally to regret her act, all things considered. If she loves me, her actual child, this is enough to block the claim that she is irrational if she does not have such regret.<sup>10</sup> Even when it implies a claim like (3), I conclude that we can accept Q.

Though Q is plausible, it does not solve the Non-Identity Problem. Q covers only the cases where, in the different outcomes, the same number of people would ever live. We need a claim that covers cases where, in the different outcomes, different numbers would ever live. The Non-Identity Problem can arise in these cases.

Because Q is restricted, it could be justified in several different ways. There are several principles that imply Q, but conflict when applied to Different Number Choices. We shall need to decide which of these principles, or which set of principles, we ought to accept. Call what we ought to accept *Theory X*. X will solve the Non-Identity Problem in Different Number Choices. And X will tell us how Q should be justified, or more fully explained.


In the case of The 14-Year-Old Girl, we are not forced to appeal to Q. There are other facts to which we could appeal, such as the effects on other people. But the problem can arise in a purer form.

### 123. HOW LOWERING THE QUALITY OF LIFE MIGHT BE WORSE FOR NO ONE

Suppose that we are choosing between two social or economic policies. And suppose that, on one of the two policies, the standard of living would be slightly higher over the next century. This effect implies another. It is not true that, whichever policy we choose, the same particular people will exist in the further future. Given the effects of two such policies on the details of our lives, it would increasingly over time be true that, on the different policies, people married different people. And, even in the same marriages, the children would increasingly over time be conceived at different times. As I have argued, children conceived more than a month earlier or later would in fact be different children. Since the choice between our two policies would affect the timing of later conceptions, some of the people who are later born would owe their existence to our choice of one of the two policies. If we had chosen the other policy, these particular people would never have existed. And the proportion of those later born who owe their existence to our choice would, like ripples in a pool, steadily grow. We can plausibly assume that, after three centuries, there would be no one living in our community who would have been born whichever policy we chose. (It may help to think about this question: how many of us could truly claim, 'Even if railways and motor cars had never been invented, I would still have been born'?)

How does this produce a problem? Consider

*Depletion.* As a community, we must choose whether to deplete or conserve certain kinds of resources. If we choose Depletion, the quality of life over the next three centuries would be slightly higher than it would have been if we had chosen Conservation. But it would later, for many centuries, be much lower than it would have been if we had chosen Conservation. This would be because, at the start of this period, people would have to find alternatives for the resources that we had depleted. It is worth distinguishing two versions of this case. The effects of the different policies would be as shown below.



We could never know, in such detail, that these would be the effects of two policies. But this is no objection to this case. Similar effects would sometimes be predictable. Nor does it matter that this imagined case is artificially simple. The case raises the questions which arise in actual cases.

Suppose that we choose Depletion, and that this has either of the two effects shown in my diagram. Is our choice worse for anyone?

Because we chose Depletion, millions of people have, for several centuries, a much lower quality of life. This quality of life is much lower, not than it is now, but than it would have been if we had chosen Conservation. These people's lives are worth living; and, if we had chosen Conservation,

these particular people would never have existed. Suppose that we do not believe that causing to exist can benefit. We should ask, 'If particular people live lives that are worth living, is this worse for these people than if they had never existed?' Our answer must be No. Suppose next that we believe that causing to exist can benefit. Since these future people's lives will be worth living, and they would never have existed if we had chosen Conservation, our choice of Depletion is not only not worse for these people: it *benefits* them.

On both answers, our choice will not be worse for these future people. Moreover, when we understand the case, we know that this is true. We know that, even if it greatly lowers the quality of life for several centuries, our choice will not be worse for anyone who ever lives.

Does this make a moral difference? There are three views. It might make all the difference, or some difference, or no difference. There might be no objection to our choice, or some objection, or the objection may be just as strong.

Some believe that *what is bad must be bad for someone*. On this view, there is no objection to our choice. Since it will be bad for no one, our choice cannot have a bad effect. The great lowering of the quality of life provides no moral reason not to choose Depletion.

Certain writers accept this conclusion.<sup>11</sup> But it is very implausible. Before we consider cases of this kind, we may accept the view that what is bad must be bad for someone. But the case of Depletion shows, I believe, that we must reject this view. The great lowering of the quality of life must provide *some* moral reason not to choose Depletion. This is believed by most of those who consider cases of this kind.

If this is what we believe, we should ask two questions:

- (1) What is the moral reason not to choose Depletion?
- (2) Does it make a moral difference that this lowering of the quality of life will be worse for no one? Would this effect be *worse*, having greater moral weight, if it *was* worse for particular people?

Our need to answer (1), and other similar questions, I call the Non-Identity Problem. This problem arises because the identities of people in the further future can be very easily affected. Some people believe that this problem is a mere quibble. This reaction is unjustified. The problem arises because of superficial facts about our reproductive system. But, though it arises in a superficial way, it is a real problem. When we are choosing between two social or economic policies, of the kind that I described, it is *not true* that, in the further future, the same people will exist whatever we choose. It is therefore *not true* that a choice like Depletion will be against the interests of future people. We cannot dismiss this problem with the pretence that this *is* true.

We partly answer question (1) if we appeal to Q. On this claim, if the

numbers would be the same, it would be bad if those who live have a lower quality of life than those who would have lived. But the problem can arise in cases where, in the different outcomes, there would be different numbers of people. To cover these cases we need Theory X. Only X will explain how Q should be justified, and provide a full solution to our problem.

#### 124. WHY AN APPEAL TO RIGHTS CANNOT WHOLLY SOLVE THE PROBLEM

Can we solve our problem by appealing to people's rights? Reconsider the 14-Year-Old Girl. By having her child so young, she gives him a bad start in life. It might be claimed: 'The objection to this girl's decision is that she violates her child's right to a good start in life'.

Even if this child has this right, it could not have been fulfilled. This girl could not have had *this* child when she was a mature woman. Some would claim that, since this child's right could not be fulfilled, this girl cannot be claimed to violate his right. The objector might reply: 'It is wrong to cause someone to exist if we know that this person will have a right that cannot be fulfilled.' Can this be the objection to this girl's decision?<sup>13</sup>

Some years ago, a British politician welcomed the fact that, in the previous year, there had been fewer teenage pregnancies. A middle-aged man wrote in anger to *The Times*. He had been born when his mother was only 14. He admitted that, because his mother was so young, his early years had been hard for both of them. But his life was now well worth living. Was the politician suggesting that it would have been better if he had never been born? This suggestion seemed to him outrageous.

The politician was, implicitly, suggesting this. On the politician's view, it would have been better if this man's mother had waited for several years before having children. I believe that we should accept this view. But can we plausibly explain this view by claiming that this angry man had a right that was not fulfilled?

I believe that we cannot. Suppose that I have a right to privacy. I ask you to marry me. If you accept, you are not acting wrongly, by violating my right to privacy. Since I am glad that you act as you do, with respect to you I *waive* this right. A similar claim applies to the writer of the angry letter to *The Times*. On the suggestion made above, this man has a right to be born by a mature woman, who would give him a good start in life. This man's mother acted wrongly because she caused him to exist with a right that cannot be fulfilled. But this man's letter shows that he was glad to be alive. He denies that his mother acted wrongly because of what she did to him. If we had claimed that her act was wrong, because he has a right that cannot be fulfilled, he could have said, 'I waive this right'. This would have undermined our objection to his mother's act.

It would have been better if this man's mother had waited. But this is not because of what she did to her actual child. It is because of what she could

have done for any child that she could have had when she was mature. The objection must be that, if she had waited, she could have given to some other child a better start in life.

Return now to the Case of Depletion. Suppose that we choose Greater Depletion. More than three centuries later, the quality of life is much lower than it would have been if we had chosen Conservation. But the people who will then be living will have a quality of life that is about as high as ours will on average be over the next century. Do these people have rights to which an objector can appeal?

It might be claimed that these people have a right to their share of the resources that we have depleted. But people do not have rights to a share of a particular resource. Suppose that we deplete some resource, but invent technology that will enable our successors, though they lack this resource, to have the same range of opportunities. There would be no objection to what we have done. The most that could be claimed is that people in each generation have a right to an equal range of opportunities, or to an equally high quality of life.<sup>14</sup>

If we choose Greater Depletion, those who live more than three centuries later will have fewer opportunities, and a lower quality of life, than some earlier and some later generations. If people have a right to equal opportunities, and an equally high quality of life, an appeal to these rights may provide some objection to our choice. Those who live more than three centuries later could not possibly have had greater opportunities, or a higher quality of life. If we had chosen otherwise, these people would never have existed. Since their rights could not be fulfilled, we may not violate their rights. But, as before, it may be objected that we cause people to exist with rights that cannot be fulfilled.

It is not clear that this is a good objection. If these people knew the facts, they would not regret that we acted as we did. If they are glad to be alive, they might react like the man who wrote to *The Times*. They might waive their rights. But I cannot assume that this is how they would all react. The appeal to these rights may provide some objection to our choice.

Can this appeal provide an objection to our choice of *Lesser Depletion*? In this case, those who live more than three centuries later have a much higher quality of life than we do now. Can we claim that these people have a *right* to an *even higher* quality of life? I believe that, on any plausible theory about rights, the answer would be No.

It will help to imagine away the Non-Identity Problem. Suppose that our reproductive system had been very different. Suppose that, whatever policies we adopt, the very same people would live more than three centuries later. The objection to our choice would then be that, for the sake of small benefits to ourselves and our children, we prevent many future people from receiving very much greater benefits. Since these future people will be better off than us, we are not acting unjustly. The objection to our act must appeal to the Principle of Utility.

Could this objection appeal to rights? Only if, like Godwin, we present Utilitarianism as a theory about rights. On Godwin's view, everyone has a right to get what the Principle of Utility implies that he should be given.<sup>15</sup> Most of those who believe in rights would reject this view. Many people explain rights as what *constrain*, or *limit*, the Principle of Utility. These people claim that it is wrong to violate certain rights, even if this would greatly increase the net sum of benefits minus burdens. On such a theory, some weight is given to the Principle of Utility. Since such a theory is not Utilitarian, this principle is better called the *Principle of Beneficence*. This principle is one part of such a theory, and the claim that we have certain rights is a different part of this theory. I shall assume that, if we believe in rights, this is the kind of moral theory that we accept.

Return to the case where we imagine away the Non-Identity Problem. If we reject Godwin's view, we could not object to the choice of Lesser Depletion by appealing to the rights of those who will live in the further future. Our objection would appeal to the Principle of Beneficence. The objection would be that, for the sake of small benefits to ourselves and our children, we deny, to people better off than us, very much greater benefits. In calling this an objection, I need not claim that it shows our choice to be wrong. I am merely claiming that, since we deny these people very much greater benefits, this provides *some* moral reason not to make this choice.

If we now restore our actual reproductive system, this reason disappears. Consider the people who will live more than three centuries later. Our choice of Lesser Depletion does not deny these people *any* benefit. If we had chosen Conservation, this would not have benefited these people, since they would never have existed.

When we assume away the Non-Identity Problem, our reason not to make this choice is explained by an appeal, not to people's rights, but to the Principle of Beneficence. When we restore the Non-Identity Problem, this reason disappears. Since this reason appealed to the Principle of Beneficence, what the problem shows is that this principle is inadequate, and must be revised. We need a better account of beneficence, or what I call Theory X.

One part of our moral theory appeals to beneficence; another part appeals to people's rights. We should therefore not expect that an appeal to rights could fill the gap in our inadequate Principle of Beneficence. We should expect that, as I have claimed, appealing to rights cannot wholly solve the Non-Identity Problem.

#### 125. DOES THE FACT OF NON-IDENTITY MAKE A MORAL DIFFERENCE?

In trying to revise our Principle of Beneficence—trying to find Theory X—we must consider cases where, in the different outcomes, different numbers of people would exist. Before we turn to these cases, we can ask



what we believe about the other question that I mentioned. Our choice of Depletion will be worse for no one. Does this make a moral difference?

We may be able to remember a time when we were concerned about effects on future generations, but had overlooked the Non-Identity Problem. We may have thought that a policy like Depletion would be against the interests of future people. When we saw that this was false, did we become less concerned about effects on future generations?

When I saw the problem, I did not become less concerned. And the same is true of many other people. I shall say that we accept the *No-Difference View*.

It is worth considering a different example:

*The Medical Programmes.* There are two rare conditions, *J* and *K*, which cannot be detected without special tests. If a pregnant woman has Condition *J*, this will cause the child she is carrying to have a certain handicap. A simple treatment would prevent this effect. If a woman has Condition *K* when she conceives a child, this will cause this child to have the same particular handicap. Condition *K* cannot be treated, but always disappears within two months. Suppose next that we have planned two medical programmes, but there are funds for only one; so one must be cancelled. In the first programme, millions of women would be tested during pregnancy. Those found to have Condition *J* would be treated. In the second programme, millions of women would be tested when they intend to try to become pregnant. Those found to have Condition *K* would be warned to postpone conception for at least two months, after which this incurable condition will have disappeared. Suppose finally that we can predict that these two programmes would achieve results in as many cases. If there is Pregnancy Testing, 1,000 children a year would be born normal rather than handicapped. If there is Preconception Testing, there would each year be born 1,000 normal children rather than a 1,000, different, handicapped children.

Would these two programmes be equally worthwhile? Let us note carefully what the difference is. As a result of either programme, 1,000 couples a year would have a normal rather than a handicapped child. These would be different couples, on the two programmes. But since the numbers would be the same, the effects on the parents and on other people would be morally equivalent. If there is a moral difference, this can only be in the effects on the children.

Note next that, in judging these effects, we need have no view about the moral status of a foetus. We can suppose that it would take a year before either kind of testing could begin. When we choose between the two programmes, none of the children has yet been conceived. And all those who are conceived will become adults. We are therefore considering effects,

not on present foetuses, but on future people. Assume next that the handicap in question, though it is not trivial, is not so severe as to make life doubtfully worth living. Even if it can be against our interests to have been born, this is not true of those born with this handicap.

Since we cannot afford both programmes, which should we cancel? Under one description, both would have the same effect. Suppose that Conditions J and K are the only causes of this handicap. The incidence is now 2,000 among those born in each year. Either programme would halve the incidence; the rate would drop to 1,000 a year. The difference is this. If we decide to cancel Pregnancy Testing, it will be true of those who are later born handicapped that, but for our decision, they would have been cured. Our decision will be worse for all these people. If instead we decide to cancel Pre-Conception Testing, there will later be just as many people who are born with this handicap. But it would not be true of these people that, but for our decision, they would have been cured. These people owe their existence to our decision. If we had not decided to cancel Pre-Conception Testing, the parents of these handicapped children would not have had *them*. They would have later had different children. Since the lives of these handicapped children are worth living, our decision will not be worse for any of them.

Does this make a moral difference? Or are the two programmes equally worthwhile? Is all that matters morally how many future lives will be lived by normal rather than handicapped people? Or does it also matter whether these lives would be lived by the very same people?

We should add one detail to the case. If we decide to cancel Pregnancy Testing, those who are later born handicapped might know that, if we had made a different decision, they would have been cured. Such knowledge might make their handicap harder to bear. We should therefore assume that, though it is not deliberately concealed, these people would not know this fact.

With this detail added, I judge the two programmes to be equally worthwhile. I know of some people who do not accept this claim, but I know of more who do.

My reaction is not merely an intuition. It is the judgement that I reach by reasoning as follows. Whichever programme is cancelled, there will later be just as many people with this handicap. These people would be different in the two outcomes that depend on our decision. And there is a claim that applies to only one of these two groups of handicapped people. Though they do not know this fact, the people in one group could have been cured. I therefore ask: 'If there will be people with some handicap, the fact that they are handicapped is bad. Would it be *worse* if, unknown to them, their handicap could have been cured?' This would be worse if this fact made these people worse off than the people whose handicap could *not* have been cured. But this fact does not have this effect. If we decide to cancel Pregnancy Testing, there will be a group of handicapped people. If we

decide to cancel Pre-Conception Testing, there will be a different group of handicapped people. The people in the first group would not be worse off than the people in the second group would have been. Since this is so, I judge these two outcomes to be morally equivalent. Given the details of the case, it seems to me irrelevant that one of the groups but not the other could have been cured.

This fact *would* have been relevant if curing this group would have reduced the incidence of this handicap. But, since we have funds for only one programme, this is not true. If we choose to cure the first group, there would be later be just as many people with this handicap. Since curing the first group would not reduce the number who will be handicapped, we ought to choose to cure this group only if they have a stronger claim to be cured. And they do not have a stronger claim. If we *could* cure the second group, they would have an equal claim to be cured. If we choose to cure the first group, they are merely luckier than the second group. Since they would merely be luckier, and they do not have a stronger claim to be cured, I do not believe that we ought to choose to cure these people. Since it is also true that, if we choose to cure these people, this will not reduce the number of people who will be handicapped, I conclude that the two programmes are equally worthwhile. If Pre-Conception Testing would achieve results in a few more cases, I would judge it to be the better programme.<sup>16</sup>

This matches my reaction to our choice of Depletion. I believe that it would be bad if there would later be a great lowering of the quality of life. And I believe that it would not be *worse* if the people who later live would themselves have existed if we had chosen Conservation. The bad effect would not be worse if it had been, in this way, worse for any particular people. In considering both cases, I accept the No-Difference View. So do many other people.

I have described two cases in which I, and many others, accept the No-Difference View. If we are right to accept this view, this may have important theoretical implications. This depends on whether we believe that, if we cause someone to exist who will have a life worth living, we thereby benefit this person. If we believe this, I cannot yet state the implications of the No-Difference View, since these will depend on decisions that I have not yet discussed. But suppose we believe that causing someone to exist cannot benefit this person. If this is what we believe, and we accept the No-Difference View, the implications are as follows.

I have suggested that we should appeal to

Q: If in either of two outcomes the same number of people would ever live, it would be bad if those who live are worse off, or have a lower quality of life, than those who would have lived.

Consider next

*The Person-Affecting View*, or *V*: It is bad if people are affected for the worse.

In Same People Choices, *Q* and *V* coincide. When we are considering these choices, those who live are the same in both outcomes. If these people are worse off, or have a lower quality of life, they are affected for the worse, and vice versa.<sup>17</sup> Since *Q* and *V* here coincide, it will make no difference to which we appeal.

The two claims conflict only in Same Number Choices. These are what this chapter has discussed. Suppose that we accept the No-Difference View. In considering these choices, we shall then appeal to *Q* rather than *V*. If we choose Depletion, this will lower the quality of life in the further future. According to *Q*, our choice has a bad effect. But, because of the facts about identity, our choice will be bad for no one. *V* does not imply that our choice has a bad effect. Would this effect be worse if it *was* worse for particular people? According to *V*, the answer must be Yes. Since we believe the No-Difference View, we answer No. We believe that *V* gives the wrong answer here. And *V* gives the wrong answer in the case of the Medical Programmes. *Q* describes the effects that we believe to be bad. And we believe that it makes no moral difference whether these effects are also bad according to *V*. *V* draws moral distinctions where, on our view, no distinctions should be drawn.

In Same People Choices, *Q* and *V* coincide. In Same Number Choices, where these claims conflict, we accept *Q* rather than *V*. When we make these two kinds of choice, we shall therefore have no use for *V*.

There remain the Different Number Choices, which *Q* does not cover. We shall here need Theory *X*. I have not yet discussed what *X* should claim. But we can predict the following. *X* will imply *Q* in Same Number Choices.

We can also predict that *X* will have the same relation to *V*. In Same People Choices, *X* and *V* will coincide. It will here make no difference to which we appeal. These are the choices with which most of our moral thinking is concerned. This explains the plausibility of *V*. This part of morality, the part concerned with beneficence, or human well-being, is usually thought of in what I shall call *person-affecting* terms. We appeal to people's interests—to what is good or bad for those people whom our acts affect. Even after we have found Theory *X*, we might continue to appeal to *V* in most cases, merely because it is more familiar. But in some cases *X* and *V* will conflict. They may conflict when we are making Same and Different Number Choices. And, whenever *X* and *V* conflict, we shall appeal to *X* rather than *V*. We shall believe that, if some effect is bad according to *X*, it makes no moral difference whether it is also bad according to *V*. As before, *V* draws a moral distinction where, on our view, no distinction should be drawn. *V* is like the claim that it is wrong to enslave whites, or to deny the vote to adult males. We shall thus conclude that this part of morality, the part concerned with beneficence and human well-being, cannot be explained

in person-affecting terms. Its fundamental principles will not be concerned with whether our acts will be good or bad for those people whom they affect. Theory X will imply that an effect is bad if it is bad for people. But this will not be *why* this effect is bad.

Remember next that these claims assume that causing to exist cannot benefit. This assumption is defensible. If we make this assumption, these claims show that many moral theories need to be revised. On these theories, it must make a moral difference whether our acts are good or bad for those people whom they affect.<sup>18</sup> And we may need to revise our beliefs about certain common cases. One example might be abortion. But most of our moral thinking would be unchanged. Many significant relations hold only between particular people. These include our relations to those to whom we have made promises, or owe gratitude, our parents, pupils, patients, clients, and (if we are politicians) those whom we represent. My remarks do not apply to such relations, or to the special obligations which they produce. My remarks apply only to our Principle of Beneficence, or our general moral reason to benefit other people, and to protect them from harm.

Since my remarks apply only to this principle, and we shall have changed our view only in some cases, this change of view may seem unimportant. This is not so. Consider once again this (too grandiose) analogy: In ordinary cases we can accept Newton's Laws. But not in all cases. And we now accept a different theory.

#### 126. CAUSING PREDICTABLE CATASTROPHES IN THE FURTHER FUTURE

In this section, rather than pursuing the main line of my argument, I discuss a minor question. In a case like that of Depletion, we cannot wholly solve the Non-Identity Problem by an appeal to people's rights. Is this also true in a variant of the case, where our choice causes a catastrophe? Since this is a minor question, this section can be ignored, except by those who do not believe that Depletion has a bad effect. Consider

*The Risky Policy.* As a community, we must choose between two energy policies. Both would be completely safe for at least three centuries, but one would have certain risks in the further future. This policy involves the burial of nuclear waste in areas where, in the next few centuries, there is no risk of an earthquake. But since this waste will remain radio-active for thousands of years, there will be risks in the distant future. If we choose this Risky Policy, the standard of living would be somewhat higher over the next century. We do choose this policy. As a result, there is a catastrophe many centuries later. Because of geological changes to the Earth's surface, an earthquake releases radiation, which kills thousands of people. Though they are killed by this catastrophe, these people will have had lives that are worth living.

We can assume that this radiation affects only people who are not yet conceived, and that its effect is to give to these people an incurable disease that will kill them at about the age of 40. This disease will have no effects before it kills.

Our choice between these two policies will affect the details of the lives that are later lived. In the way explained above, our choice will therefore affect who will later live. After many centuries there would be no one living in our community who, whichever policy we chose, would have been born. Because we chose the Risky Policy, thousands of people are later killed. But if we had chosen the alternative Safe Policy, these particular people would never have existed. Different people would have existed in their place. Is our choice of the Risky Policy worse for anyone?

We should ask, 'If people live lives that are worth living, even though they are killed by some catastrophe, is this worse for these people than if they had never existed?' Our answer must be No. Though it causes a predictable catastrophe, our choice of the Risky Policy will be worse for no one.

Some may claim that our choice of Depletion does not have a bad effect. This cannot be claimed about our choice of the Risky Policy. Since this choice causes a catastrophe, it clearly has a bad effect. But our choice will not be bad for, or worse for, any of the people who later live. This case forces us to reject the view that a choice cannot have a bad effect if this choice will be bad for no one.

In this case, the Non-Identity Problem may seem easier to solve. Though our choice is not worse for the people struck by the catastrophe, it might be claimed that we harm these people. And the appeal to people's rights may here succeed.

We can deserve to be blamed for harming others, even when this is not worse for them. Suppose that I drive carelessly, and in the resulting crash cause you to lose a leg. One year later, war breaks out. If you had not lost this leg, you would have been conscripted, and killed. My careless driving therefore saves your life. But I am still morally to blame.

This case reminds us that, in assigning blame, we must consider not actual but predictable effects. I knew that my careless driving might harm others, but I could not know that it would in fact save your life. This distinction might apply to our choice of the Risky Policy. Suppose we know that, if we choose this policy, this may in the distant future cause many accidental deaths. But we have overlooked the Non-Identity Problem. We mistakenly believe that, whichever policy we choose, the very same people will later live. We therefore believe that our choice of the Risky Policy may be very greatly against the interests of some future people. If we believe this, our choice can be criticized. We can deserve blame for doing what we *believe* may be greatly against the interests of other people. This criticism stands even if our belief is false—just as I am as much to blame even if my

careless driving will in fact save your life.

Suppose that we cannot find Theory X, or that X seems less plausible than the objection to doing what may be greatly against the interests of other people. It may then be better if we conceal the Non-Identity Problem from those who will decide whether we increase our use of nuclear energy. It may be better if these people believe falsely that such a policy may, by causing a catastrophe, be greatly against the interests of some of those who will live in the distant future. If these people have this false belief, they may be more likely to reach the right conclusions.

We have lost this false belief. We realize that, if we choose the Risky Policy, our choice will *not* be worse for those people whom the catastrophe later kills. Note that this is not a lucky guess. It is not like predicting that, if I cause you to lose a leg, this will later save you from death in the trenches. We know that, if we choose the Risky Policy, this may in the distant future cause many people to be killed. But we also know that, if we had chosen the Safe Policy, the people who are killed would never have been born. Since these people's lives will be worth living, we *know* that our choice will not be worse for them.

If we know this, we cannot be compared to a careless driver. What is the objection to our choice? Can it be wrong to harm others, when we know that our act will not be worse for the people harmed? This might be wrong if we could have asked these people for their consent, but have failed to do so. By failing to ask these people for their consent, we infringe their autonomy. But this cannot be the objection to our choice of the Risky Policy. Since we could not possibly communicate with the people living many centuries from now, we cannot ask for their consent.

When we cannot ask for someone's consent, we should ask instead whether this person would later regret what we are doing. Would the people who are later killed regret our choice of the Risky Policy? Let us suppose that these people know all of the facts. From an early age they know that, because of the release of radiation, they have an incurable disease that will kill them at about the age of 40. They also know that, if we had chosen the Safe Policy, they would never have been born. These people would regret the fact that they will die young. But, since their lives are worth living, they would not regret the fact that they were ever born. They would therefore not regret our choice of the Risky Policy.

Can it be wrong to harm others, when we know *both* that if the people harmed knew about our act, they would not regret this act, *and* that our act will not be worse for these people than anything else that we could have done? How might we know that, though we are harming someone, our act will not be worse for this person? There are at least two kinds of case:

(1) Though we are harming someone, we may also know that we are giving to this person some fully compensating benefit. We could not know this unless the benefit would clearly outweigh the harm. But, if this is so, what we are doing will be better for this person. In this kind of case, if we

are also not infringing this person's autonomy, there is no objection to our act. There is no objection to our harming someone when we know both that this person will have no regrets, and that our act will be clearly better for this person. In English Law, surgery was once regarded as justifiable grievous bodily harm. As I argued in Section 25, we should revise the ordinary use of the word 'harm'. If what we are doing will not be worse for some other person, or will even be better for this person, we are not, in a morally relevant sense, harming this person.

If we assume that causing to exist can benefit, our choice of the Risky Policy is, in its effects on those killed, like the case of the surgeon. Though our choice causes these people to be killed, since it also causes them to exist with a life worth living, it gives them a benefit that outweighs this harm. The objection to our choice cannot be that it harms these people.

We may instead assume that causing to exist cannot benefit. On this assumption, our choice of the Risky Policy does not give to the people whom it kills some fully compensating benefit. Our choice is not *better* for these people. It is merely *not worse* for them.

(2) There is another kind of case in which we can know that, though we are harming someone on the ordinary use of 'harm', this will not be worse for this person. These are the cases that involve overdetermination. In these cases we know that, if we do not harm someone, this person will be harmed at least as much in some other way. Suppose that someone is trapped in a wreck and about to be burnt to death. This person asks us to shoot him, so that he does not die painfully. If we kill this person we are not, in a morally relevant sense, harming him.

Such a case cannot show that there is no objection to our choice of the Risky Policy, since it is not relevantly similar. If the catastrophe did not occur, the people killed would have lived for many more years. There is a quite different reason why our choice of the Risky Policy is not worse for these people.

Could there be a case in which we kill some existing person, knowing what we know when we choose the Risky Policy? We must know (*a*) that this person will learn but not regret the fact that we have done something that will cause him to be killed. And we must know (*b*) that, though this person would otherwise have lived a normal life for many more years, causing him to be killed will be neither better nor worse for him. ((*b*) is what we know about the effects of our choice of the Risky Policy, if we assume that, in doing what is a necessary part of the cause of the existence of the people killed by the catastrophe, we cannot be benefiting these people.)

Suppose that we kill some existing person, who would otherwise have lived a normal life for many more years. In such a case, we could not *know* that (*b*) is true. Even if living for these many years would be neither better nor worse for this person, this could never be predicted. There cannot be a case where we kill some existing person, knowing what we know when we choose the Risky Policy. A case that is relevantly similar must involve



causing someone to be killed who, if we had acted otherwise, would never have existed.

Compare these two cases:

*Jane's Choice.* Jane has a congenital disease, that will kill her painlessly at about the age of 40. This disease has no effects before it kills. Jane knows that, if she has a child, it will have this same disease. Suppose that she can also assume the following. Like herself, her child would have a life that is worth living. There are no children who need to be but have not been adopted. Given the size of the world's population when this case occurs (perhaps in some future century), if Jane has a child, this will not be worse for other people. And, if she does not have this child, she will be unable to raise a child. She cannot persuade someone else to have an extra child, whom she would raise. (These assumptions give us the relevant question.) Knowing these facts, Jane chooses to have a child.

*Ruth's Choice.* Ruth's situation is just like Jane's, with one exception. Her congenital disease, unlike Jane's, kills only males. If Ruth pays for the new technique of in vitro fertilization, she would be certain to have a daughter whom this disease would not kill. She decides to save this expense, and takes a risk. Unluckily, she has a son, whose inherited disease will kill him at about the age of 40.

Is there a moral objection to Jane's choice? Given the assumptions in the case, this objection would have to appeal to the effect on Jane's child. Her choice will not be worse for this child. Is there an objection to her choice that appeals to this child's rights? Suppose we believe that each person has a right to live a full life. Jane knows that, if she has a child, his right to a full life could not possibly be fulfilled. This may imply that Jane does not violate this right. But the objection could be restated. It could be said: 'It is wrong to cause someone to exist with a right that cannot be fulfilled. This is why Jane acts wrongly.'

Is this a good objection? If I was Jane's child, my view would be like that of the man who wrote to *The Times*. I would regret the fact that I shall die young. But, since my life is worth living, I would not regret that my mother caused me to exist. And I would deny that her act was wrong because of what it did to me. If I was told that it *was* wrong, because it caused me to exist with a right that cannot be fulfilled, I would *waive* this right.

If Jane's child waives his right, this undermines this objection to her choice. But, though *I* would waive this right, I cannot be certain that, in all such cases, this is what such a child would do. If Jane's child does not waive his right, the appeal to this right may perhaps provide some objection to her choice.

Turn now to Ruth's choice. There is clearly a greater objection to *this* choice. This is because Ruth has a different alternative. If Jane does not

have a child, she will not be able to raise a child; and one fewer life will be lived. Ruth's alternative is to pay for the technique that will give her a different child whom her disease will not kill. She chooses to save this expense, knowing that the chance is one in two that her child will be killed by this disease.

Even if there is an objection to Jane's choice, there is a greater objection to Ruth's choice. This objection cannot appeal only to the effects on Ruth's actual child, since these are just like the effects of Jane's choice on Jane's child. The objection to Ruth's choice must appeal in part to the possible effect on a different child who, by paying for the new technique, she could have had. The appeal to this effect is not an appeal to anyone's rights.

Return now to our choice of the Risky Policy. If we choose this policy, this may cause people to exist who will be killed in a catastrophe. We know that our choice would not be worse for these people. But, if there is force in the objection to Jane's choice, this objection would apply to our choice. By choosing the Risky Policy, we may cause people to exist whose right to a full life cannot be fulfilled.

The appeal to these people's rights may provide some objection to our choice. But it cannot provide the whole objection. Our choice is, in one respect, unlike Jane's. Her alternative was to have no child. Our alternative is like Ruth's. If we had chosen the Safe Policy, we would have had different descendants, none of whom would have been killed by released radiation.

The objection to Ruth's choice cannot appeal only to her child's right to a full life. The same is therefore true of the objection to our choice of the Risky Policy. This objection must in part appeal to the effects on the possible people who, if we had chosen differently, would have lived. As before, the appeal to rights cannot wholly solve the Non-Identity problem. We must also appeal to a claim like Q, which compares two different sets of possible lives.

It may be objected: 'When Ruth conceives her child, it inherits the disease that will deny it a full life. Because this child's disease is inherited in this way, it cannot be claimed that Ruth's choice kills her child. If we choose the Risky Policy, the causal connections are less close. Because the connections are less close, our choice kills the people who later die from the effects of released radiation. That we kill these people is the full objection to our choice.'

This objection I find dubious. Why is there a greater objection to our choice because the causal connections are less close? The objection may be correct in what it claims about our ordinary use of 'kill'. But, as I argued in Section 25, this use is morally irrelevant. Since that argument may not convince, I add

*The Risky Cure for Infertility.* Ann cannot have a child unless she takes a certain treatment. If she takes this treatment, she will have a son, who will be healthy. But there is a risk of one in two that this treatment will

give Ann a rare disease. This disease has the following features. It remains dormant for twenty years, is undetectable, kills men but does not harm women, and is infectious. The following is therefore true. If Ann takes this treatment and has a healthy son, there is a chance of one in two that in twenty years she will infect her son with a disease that will kill him twenty years later, or when he is about forty. Ann chooses to take this treatment, and she does later infect her son with this fatal disease.

On the objection stated above, there is a strong objection to Ann's choice, which does not apply to Ruth's choice. Because the causal connections are less close, Ann's choice kills her son. And she knew that the chance was one in two that her choice would have this effect. Ruth knows that there is the same chance that her child will die at about the age of 40. But, because the causal connections are so close, her choice does not kill her son. According to this objection, this difference has great moral relevance.

This is not plausible. Ruth and Ann both know that, if they act in a certain way, there is a chance of one in two that they will have sons who will be killed by a disease at about the age of forty. The causal story is different. But this does not make Ann's choice morally worse. I believe that this example shows that we should reject this last objection.

The objector might say: 'I deny that, by choosing to take the Risky Cure, Ann kills her son'. But, if the objector denies this, he cannot claim that, by choosing the Risky Policy, we kill some people in the distant future. The causal connections take the same form. Each choice produces a side-effect which later kills people who owe their existence to this choice.

If this objection fails, as I believe, my earlier claim is justified. It is morally significant that, if we choose the Risky Policy, our choice is like Ruth's rather than Jane's. It is morally significant that, if we had chosen otherwise, different people would have lived who would not have been killed. Since this is so, the objection to our choice cannot appeal only to the rights of those who actually later live. It must also appeal to a claim like Q, which compares different sets of possible lives. As I claimed earlier, the appeal to rights cannot wholly solve the Non-Identity Problem.

## 127. CONCLUSIONS

I shall now summarize what I have claimed. It is in fact true of everyone that, if he had not been conceived within a month of the time when he was conceived, he would never have existed. Because this is true, we can easily affect the identities of future people, or *who* the people are who will later live. If a choice between two social policies will affect the standard of living or the quality of life for about a century, it will affect the details of all the lives that are later lived in our community. As a result, some of those who later live will owe their existence to our choice of one of these two policies.

After three centuries, this will be true of everyone in our community.

This fact produces a problem. One of these two policies may, in the further future, cause a great lowering of the quality of life. This would be the effect of the policy I call Depletion. This effect is bad, and provides a moral reason not to choose Depletion. But, because of the fact just mentioned, our choice of Depletion will be worse for no one. Some people believe that a choice cannot have bad effects if this choice will be worse for no one. The Case of Depletion shows that we must reject this view. And this is shown more forcefully by the Case of the Risky Policy. One effect of choosing this policy is a catastrophe that kills thousands of people. This effect is clearly bad, even though our choice will be worse for no one.

Since these two choices will be worse for no one, we need to explain why we have a moral reason not to make these choices. This problem arises because, in the different outcomes, different people would exist. I therefore call this the Non-Identity Problem.

I asked whether we can solve this problem by appealing to people's rights. I argued that, even in the case of the Risky Policy, the objection to our choice cannot appeal only to people's rights. The objection must in part appeal to a claim like Q, which compares different possible lives. And we cannot plausibly appeal to rights in explaining the objection to our choice of Lesser Depletion. Even after the great lowering of the quality of life, those who will be living will be much better off than we are now. These people cannot be claimed to have a *right* to the even higher quality of life that different people would have enjoyed if we had chosen Conservation. If we imagine away the Non-Identity Problem, the objection to our choice would appeal to our Principle of Beneficence. To solve the Non-Identity Problem, we must revise this principle.

One revised principle is Q, the Same Number Quality Claim. According to Q, if in either of two outcomes there would be the same number of people, it would be bad if those who live are worse off, or have a lower quality of life, than those who would have lived. We need a wider principle to cover cases where, in the different outcomes, there would be different numbers of people. This needed principle I call Theory X. Only X will fully solve the Non-Identity Problem.

Does the fact of non-identity make a moral difference? When we see that our choice of Depletion will be worse for no one, we may believe that there is less objection to our choice. But I believe that the objection is just as strong. And I have a similar belief when I compare the effects of the two Medical Programmes. This belief I call the No Difference View. Though I know of some people who do not accept this view, I know of more who do. If we accept the No Difference View, and believe that causing to exist cannot benefit, this has wide theoretical implications. We can predict that Theory X will not take a person-affecting form. The best theory about beneficence will not appeal to what is good or bad for those people whom our acts affect.

In what follows, I shall try to find Theory X. As I have claimed, this attempt will raise some puzzling questions.